

# Cloud and Machine Learning

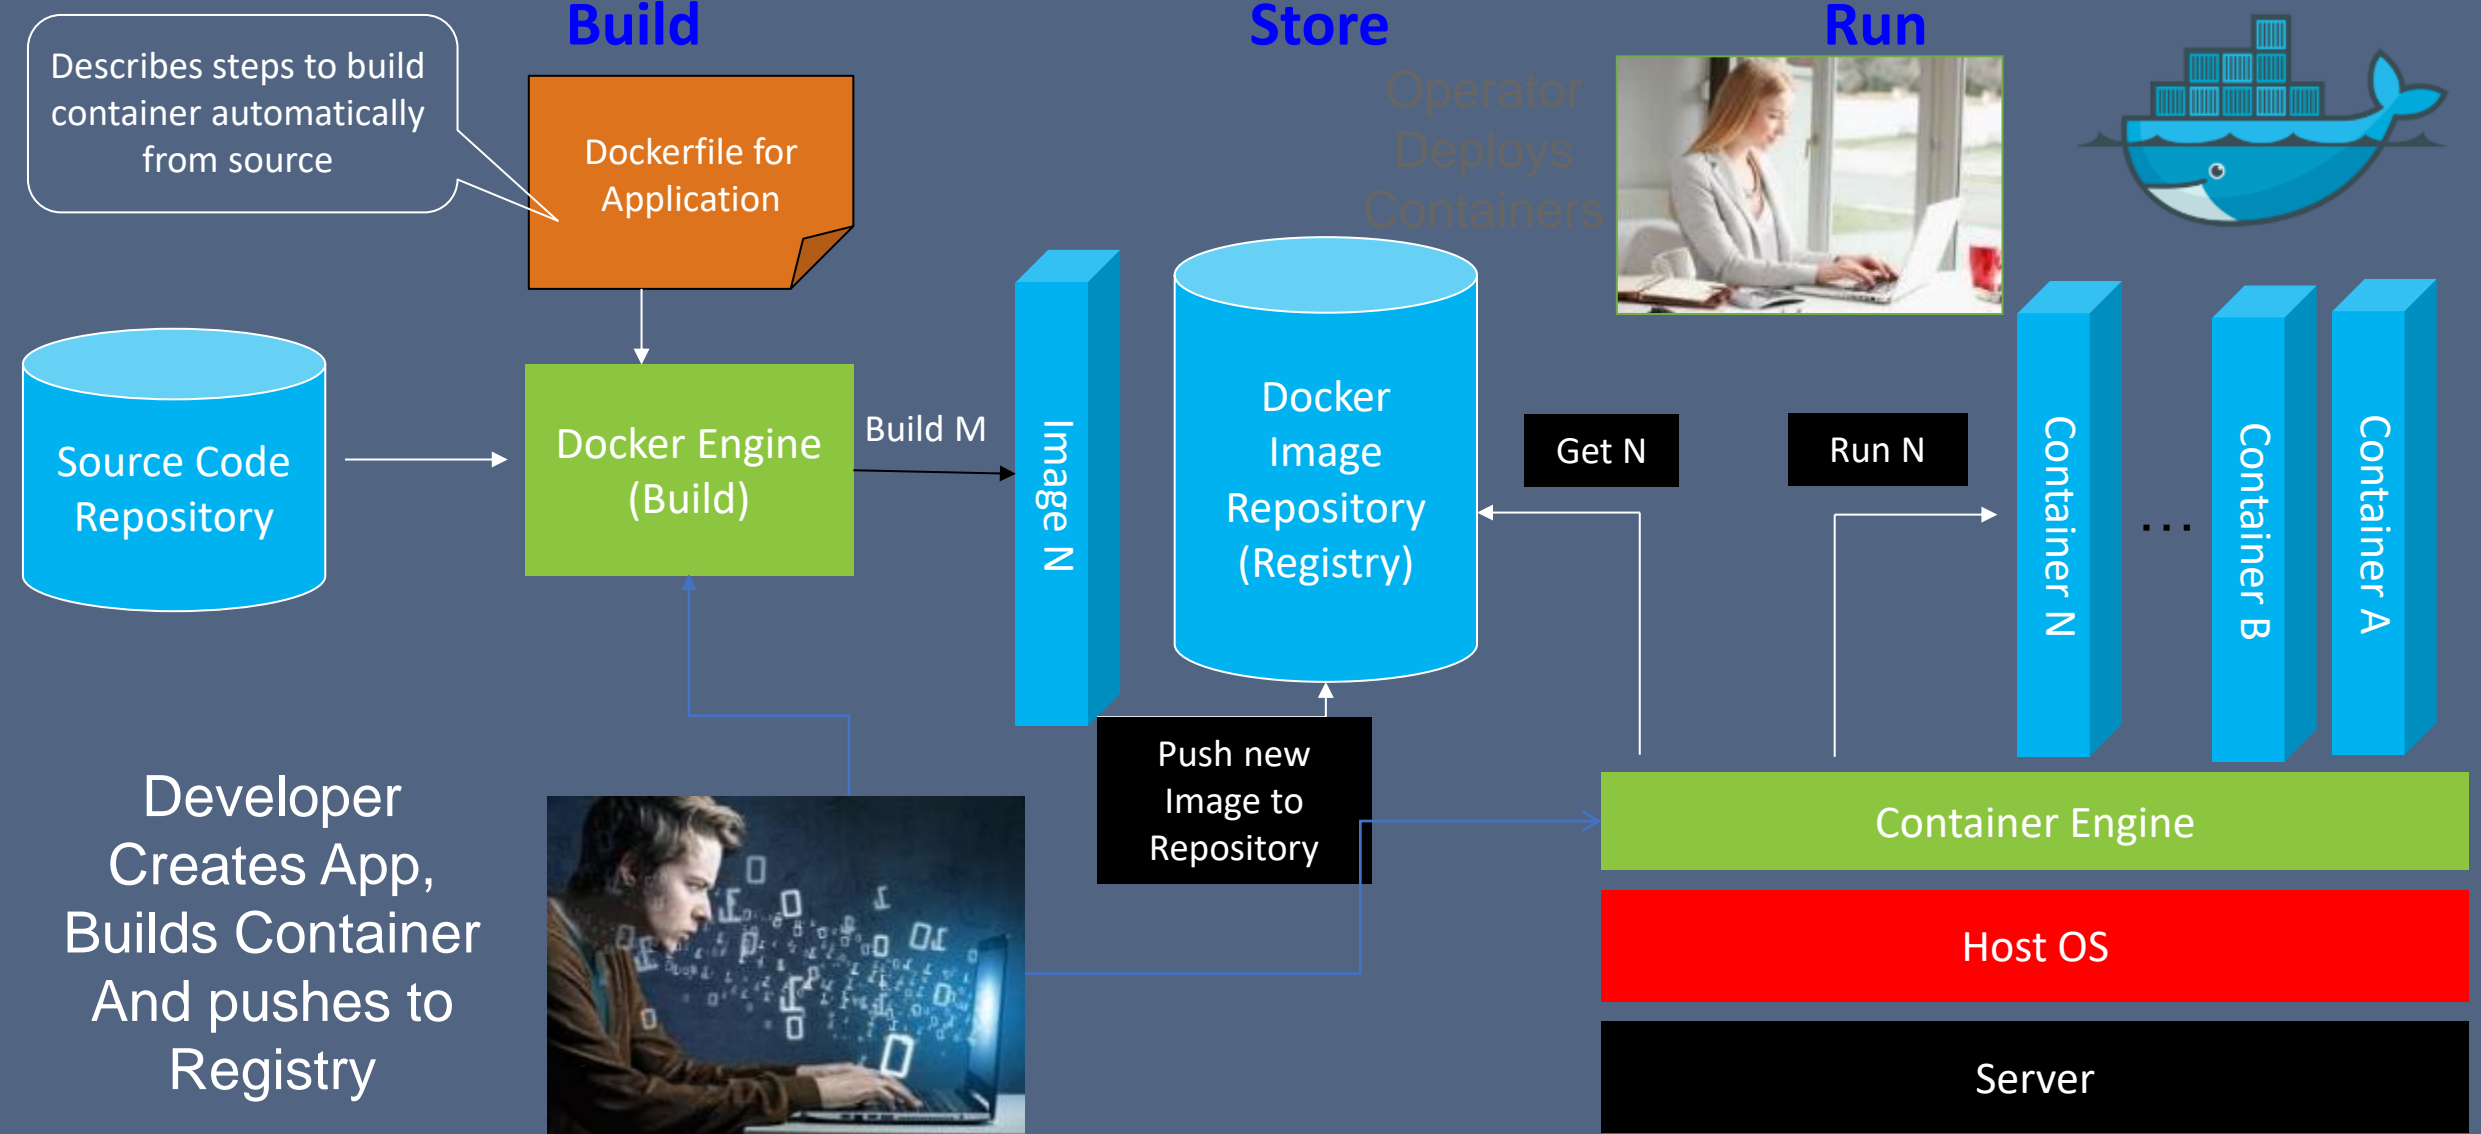
## CSCI-GA.3033-085 Fall 2024

Prof. Hao Yu

Prof. I-Hsin Chung

Lecture 8: Kubernetes

# What are the Basic Functions of Containers



# Docker MNIST HW – graded

- **Work:** Run MNIST inside a Docker container
  - MNIST can be based on PyTorch or TensorFlow.
  - We provided an example code set.
- **Submission:**
  - Source code, files containing your own work or being modified by you (35 points)
    - Vagrantfile (If need to use Vagrant for your experiment)
    - Dockerfile (must include one or a few changes to the example, i.e. mnist/main.py command line options)
    - On screen output (partial capture) of the 'docker run'.
    - Other downloaded files with you modifications
  - Report (55 points)
    - Document your steps so we can reproduce them. These should include a workflow, observations with attention to details, map to learnt concepts, code comment or explanation, and brief comparative discuss.
    - The purpose of the report is show that you can not only do the operation, but have conceptual build-up of container technologies.
  - **Singularity** (10 points):
    - Run MNIST inside a Singularity container (inside a Vagrant VM if need)
  - Compare Docker and Singularity in your report.

# Docker MNIST HW – Exemplary

- A possible work flow:

- `git clone https://github.com/yuhaohaoyu/ibmcloud-fall-2023.git`
- `git clone https://github.com/pytorch/examples.git`
- `cd ibmcloud-fall-2023; cp -r ../examples/mnist .`
- `vi Vagrantfile Dockerfile`
- `vagrant up; vagrant ssh; cd /vagrant;`
- `docker build -t mnist . ;`
- `docker run -it mnist 2>&1 | tee docker-run.out`

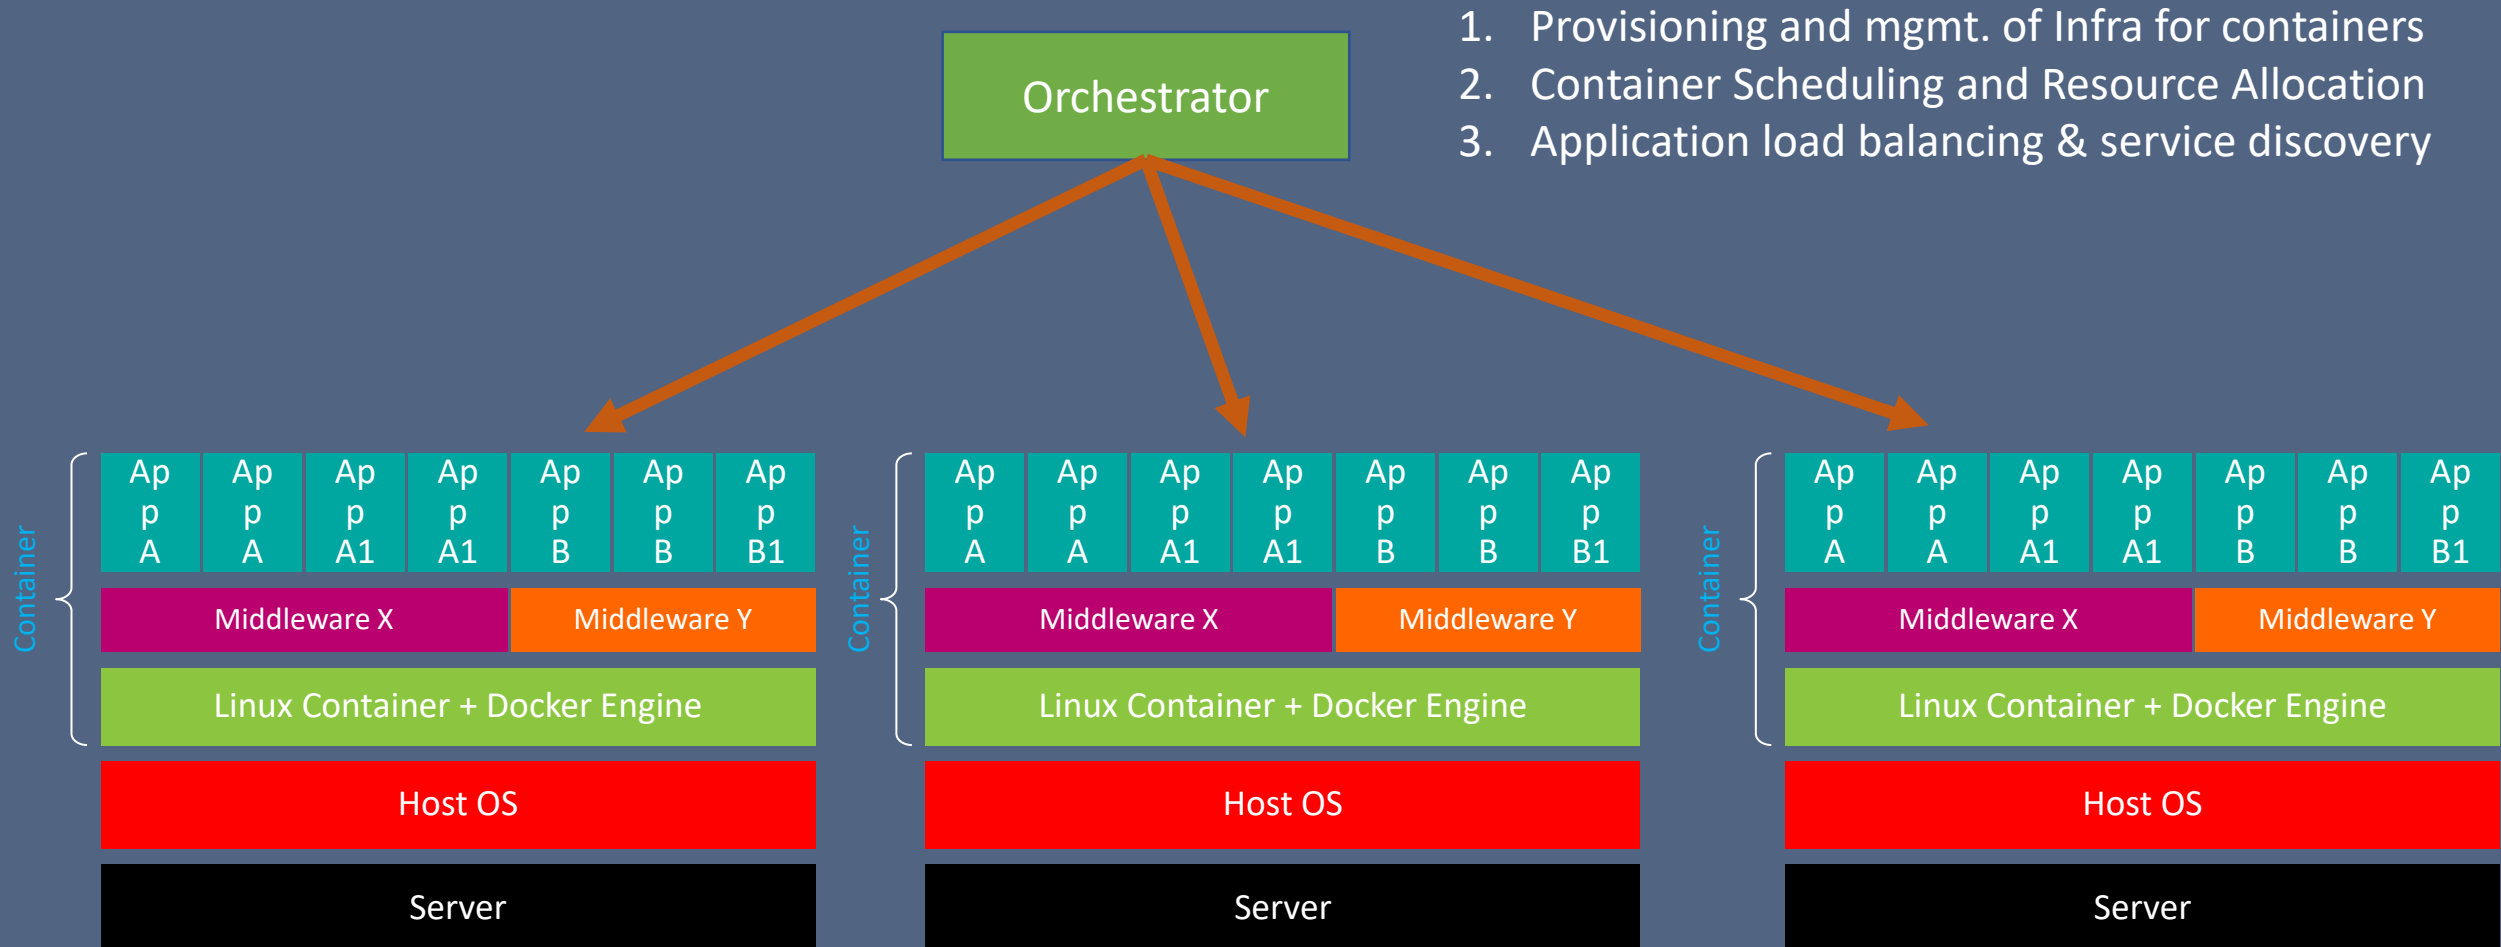
- Submission:

- Vagrantfile (only if need), Dockerfile, docker-run.out, mnist/main.py file if there is change.

- References:

- Docker hands-on slides.
- Vagrant and Docker hands-on code: <https://github.com/pytorch/examples.git>
- Install docker on laptop:
  - Install docker with administrator privilege
    - <https://www.geeksforgeeks.org/how-to-install-docker-on-windows/>
  - Install docker in a virtual machine with root privilege: e.g. Vagrant-VM
    - <https://docs.docker.com/engine/install/ubuntu/>

# What is container orchestration?

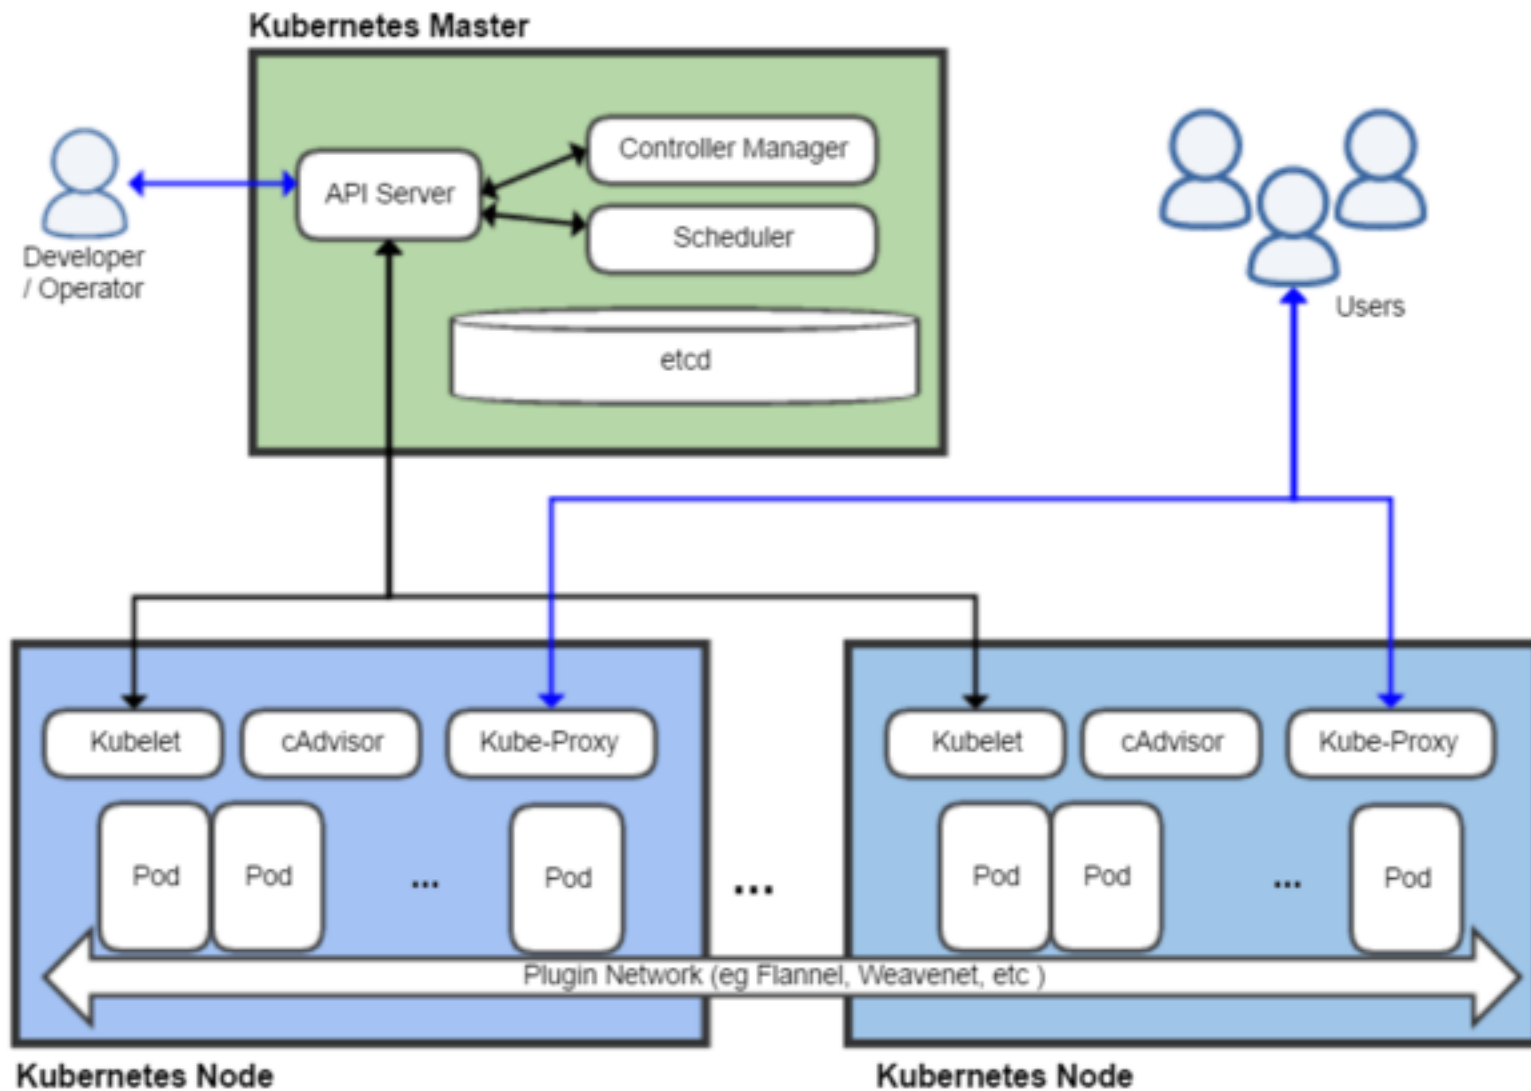


# Container orchestrators

- Apache Mesos & DCOS
- Mesos Marathon
- Docker Swarm
- Docker Compose
- CoreOS Fleet
- Tupperware (twine)
- OpenShift
- Kubernetes

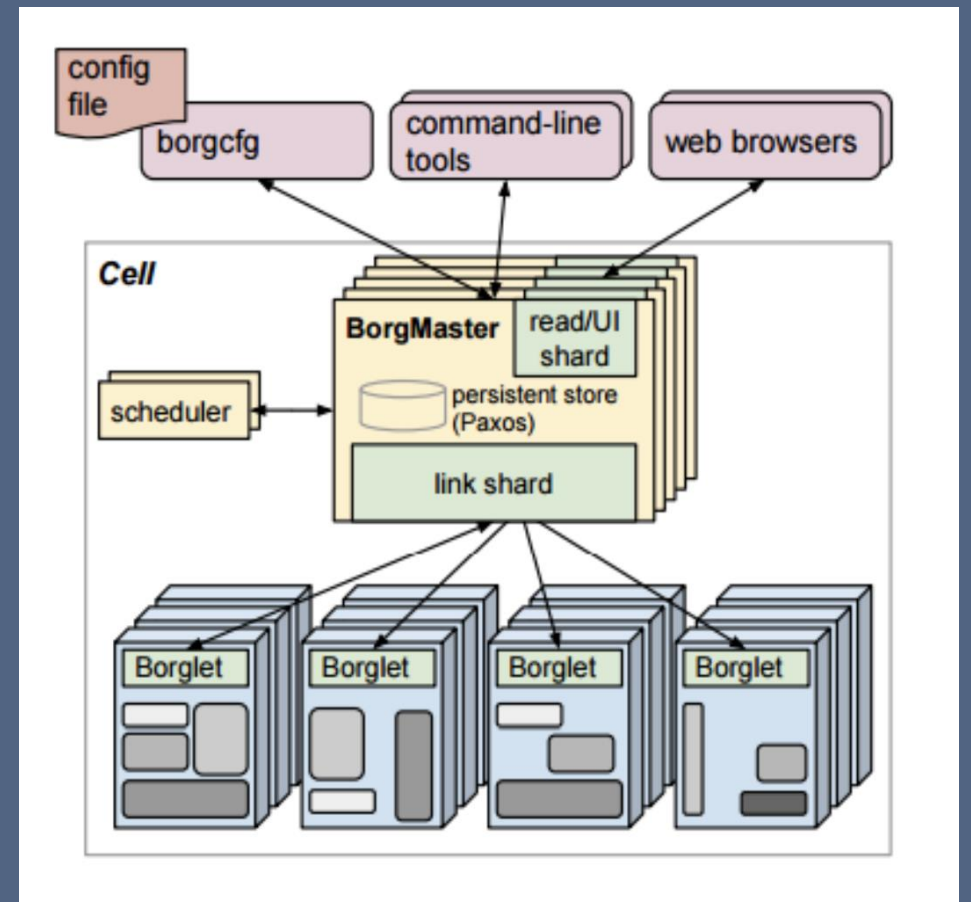
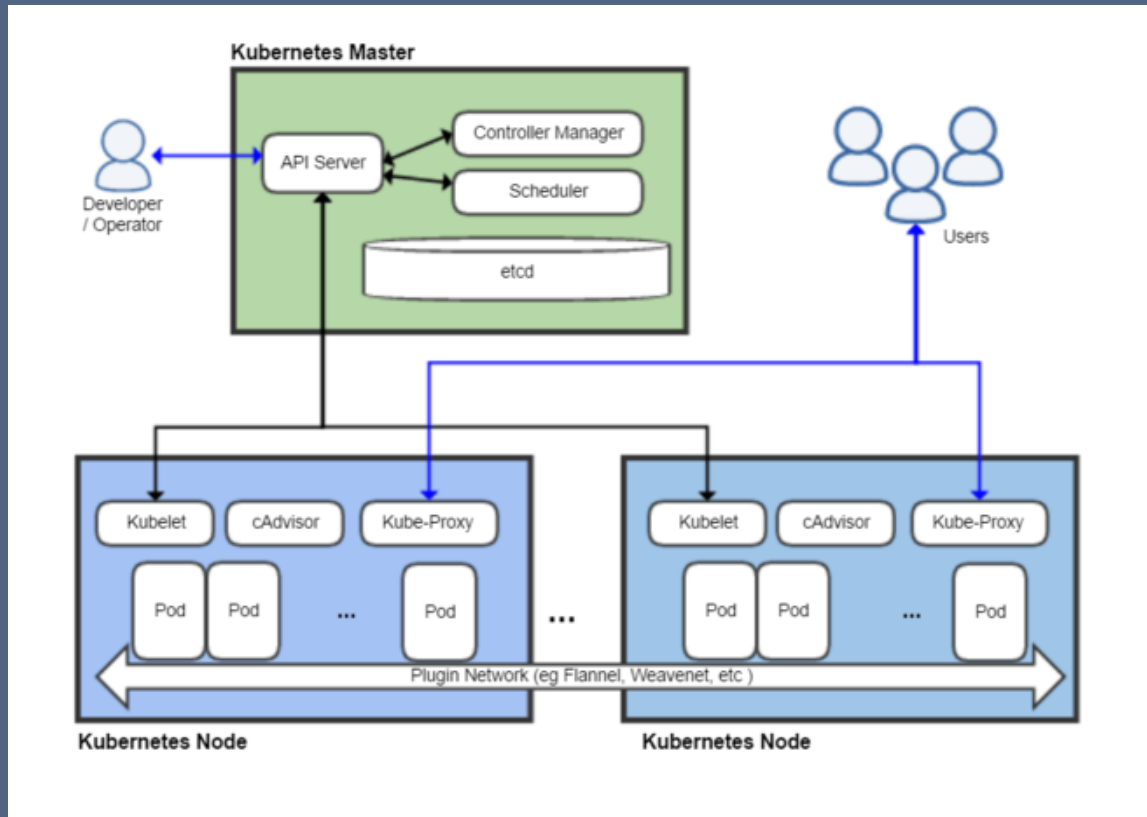
# What is Kubernetes?

- “Kubernetes is a portable, extensible open-source platform for managing containerized **workloads** and **services**, that facilitates both **declarative configuration** and **automation**. It has a large, rapidly growing ecosystem. Kubernetes services, support, and tools are widely available.”
  - From: <https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/>





# Kubernetes vs Borg



# Kubernetes basic objects and controllers

- Basic Kubernetes objects

- **Pod**
- **Service**
- **Volume**
- **Namespace**


- Kubernetes Machinery

- **Kube API Server**
- **Kube Scheduler**
- **Controller-manager**
- **Kubelet**
- **Kube-proxy**


- Controllers

- **ReplicaSet**
- **Deployment**
- **StatefulSet**
- **DaemonSet**
- **Job**

# Nodes & Pods



# Nodes & Pods




Node:

1. Kubelet (communication between master → Node)
2. Container runtime (docker, cri-o, podman)

# Nodes & Pods



# Nodes & Pods



# Nodes & Pods




# Nodes & Pods






# Nodes & Pods




# Nodes & Pods



Pod:

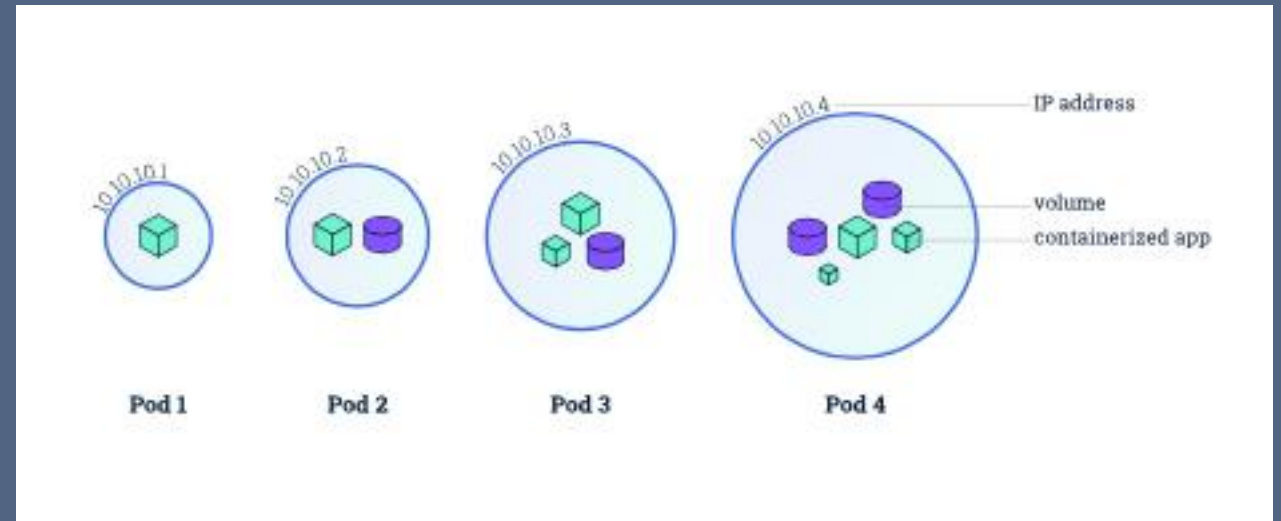
1. Containers run in Pods
2. Shared storage
3. Shared network
4. Container behavior (ports, health, resources)
5. All containers of a pod run on the same machine

# Nodes & Pods



Node:

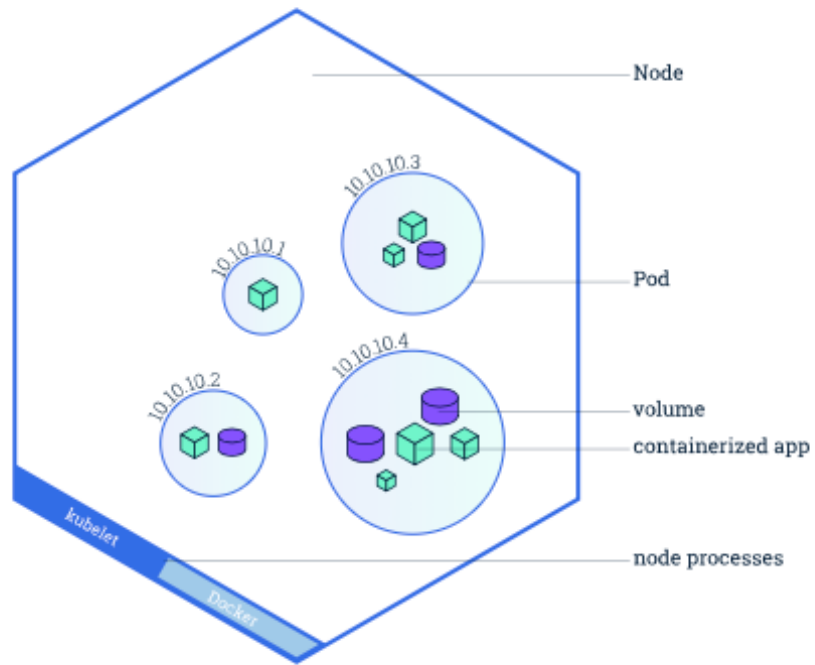
1. Kubelet (communication between master → Node)
2. Container runtime (docker, rkt)



Pod:

1. Containers run in Pods
2. Shared storage
3. Shared network
4. Container behavior (ports, health, resources)
5. All containers of a pod run on the same machine

# K8S Pod vs real PODS



# “Hello world” pod spec

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  containers:
  - name: myapp-container
    image: busybox
    command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 3600']
```

```
[~]$ kubectl create -f hello.yaml
```

```
pod/myapp-pod created
```

```
[~]$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
ibm-cert-manager-cert-manager-768b66977-qp6db	1/1	Running	0	12d
myapp-pod	0/1	ContainerCreating	0	5s
test-pd	1/1	Running	0	2d10h

```
[~]$
```


# Differences between Docker and K8S

```
docker run busybox /bin/sh -c 'echo hello && sleep 10'
```

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  containers:
  - name: myapp-container
    image: busybox
    command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 3600']
```


# Pod spec to Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  containers:
  - name: myapp-container
    image: busybox
    command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 3600']
```



# Kubernetes ReplicaSets

- ReplicaSets allow multiple copies of pods to be deployed
- One or more containers in a pod
- If a container dies, the ReplicaSet will spawn a new one





# ReplicaSet

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-rs
  labels:
    app: myapp-rs
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: myapp-rs
  template:
    metadata:
      labels:
        tier: myapp-rs
    spec:
      containers:
      - name: myapp-container
        image: busybox
        command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 3600']
```

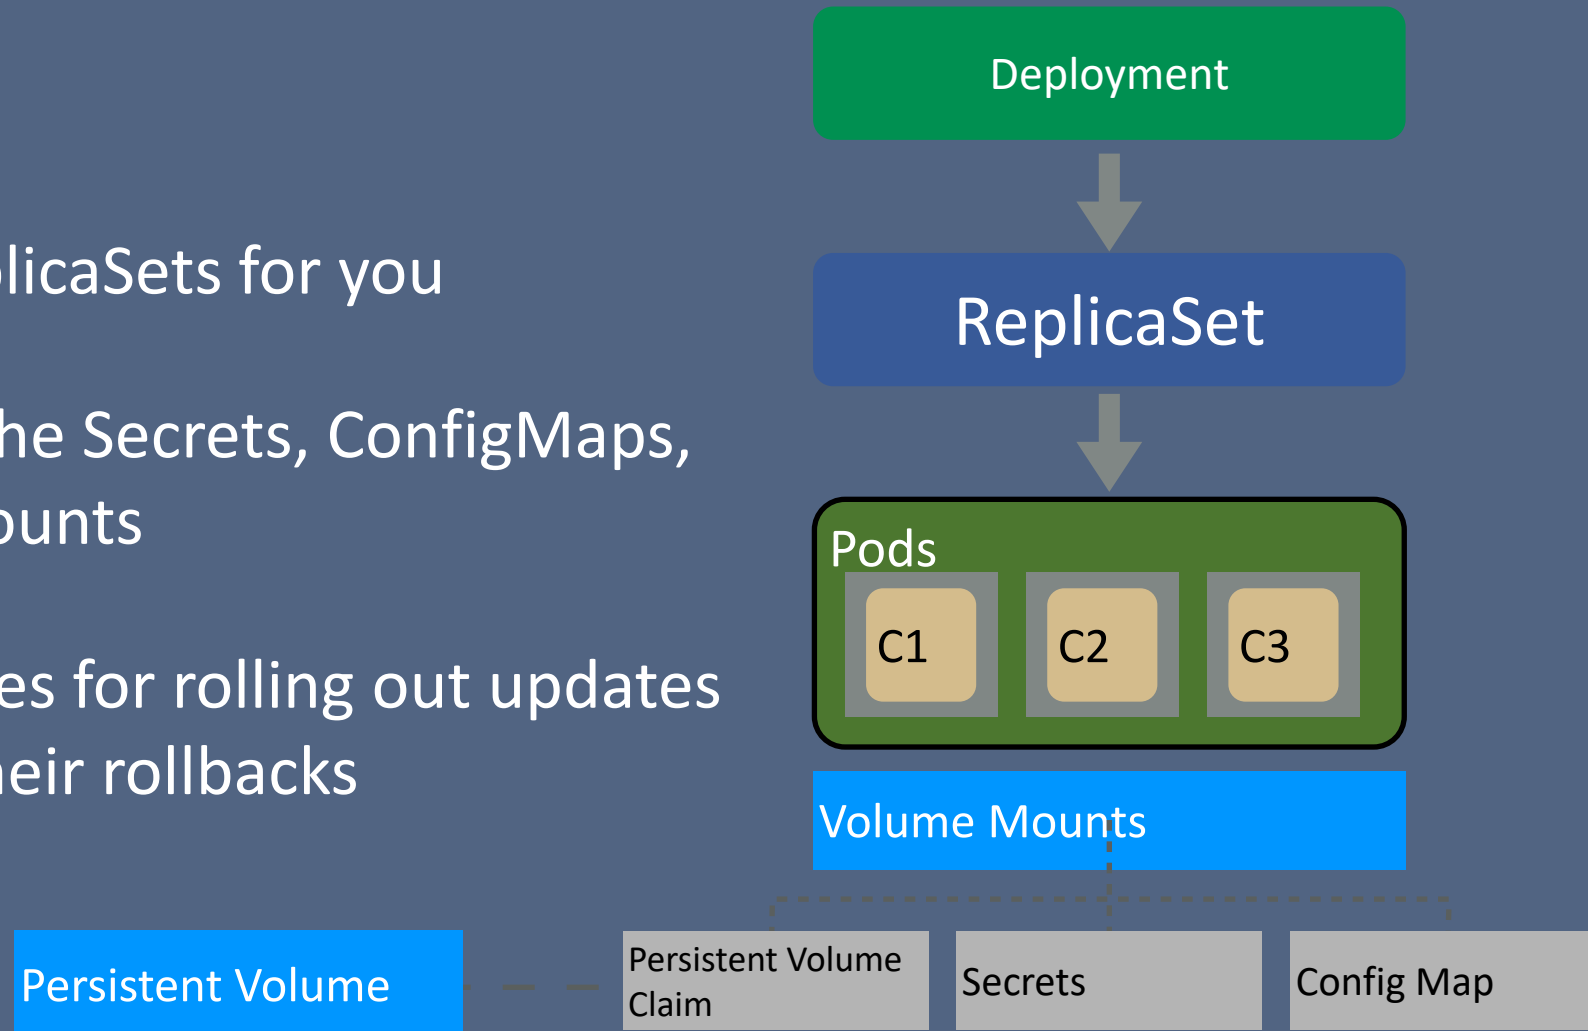
# ReplicaSet vs Pod

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-rs
  labels:
    app: myapp-rs
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: myapp-rs
  template:
    metadata:
      labels:
        tier: myapp-rs
    spec:
      containers:
      - name: myapp-container
        image: busybox
        command: ['sh', '-c', 'echo Hello Kubernetes
```

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  containers:
  - name: myapp-container
    image: busybox
    command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 36']
```

# Kubernetes Deployments

- Deployment
  - Sets up the ReplicaSets for you
  - Also specified the Secrets, ConfigMaps, and Volume Mounts
  - Provides features for rolling out updates and handling their rollbacks

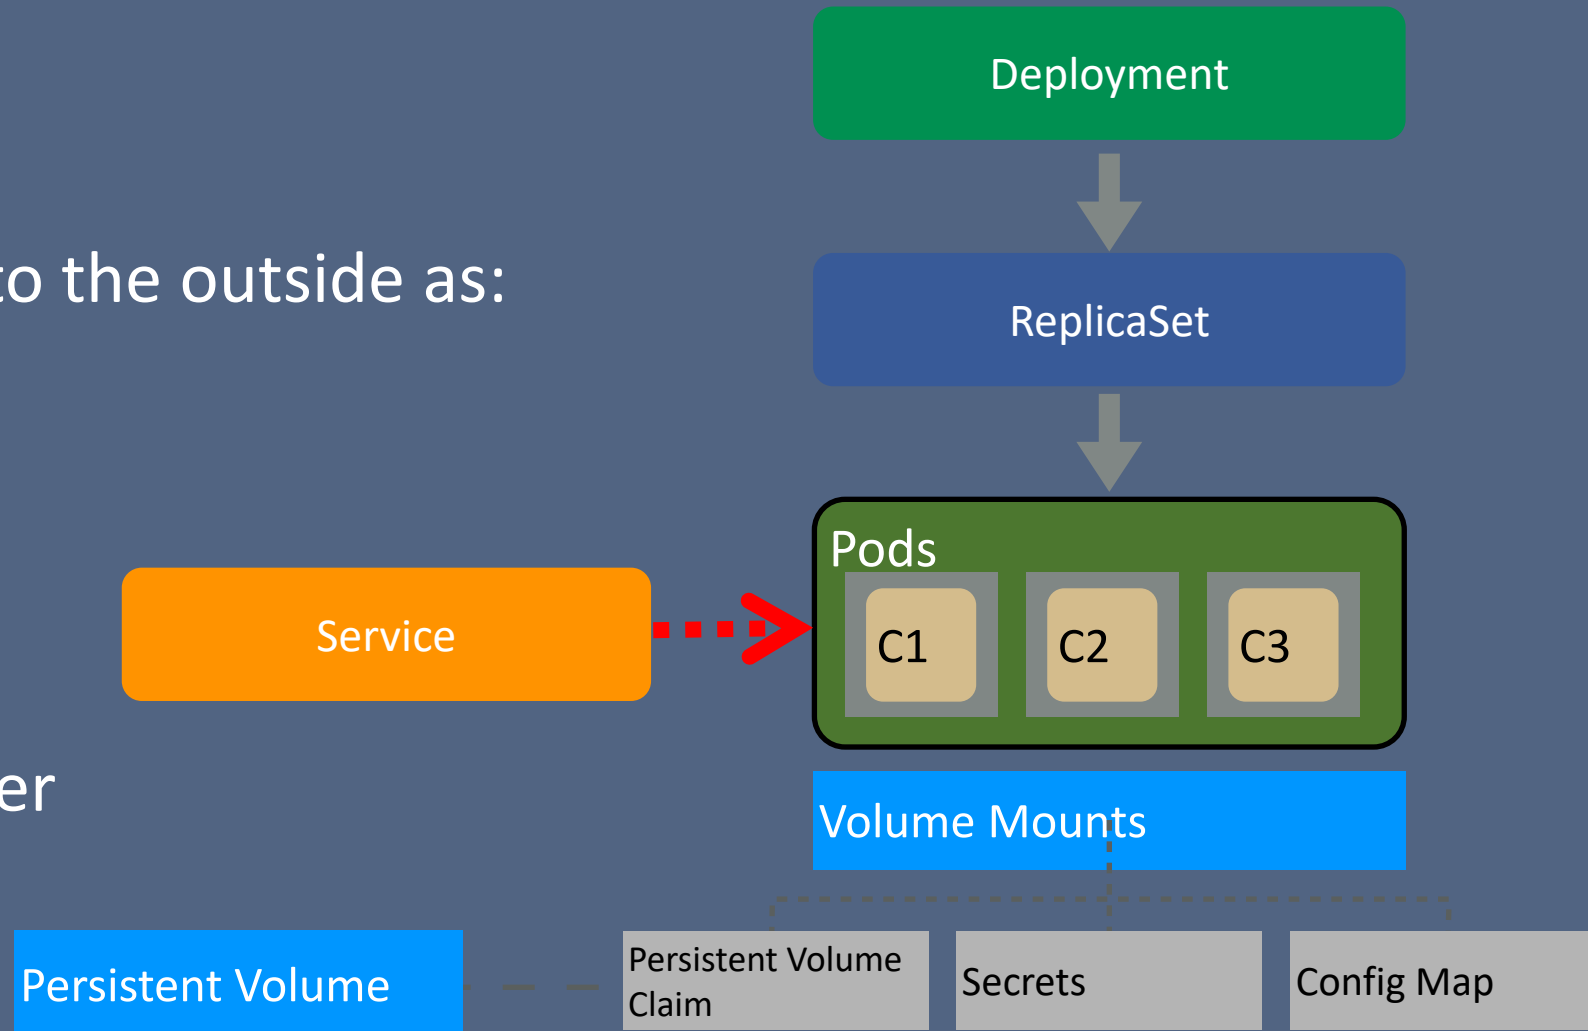


# Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-deployment
  labels:
    app: hello-d
spec:
  replicas: 3
  selector:
    matchLabels:
      app: hello-d
  template:
    metadata:
      labels:
        app: hello-d
    spec:
      spec:
        containers:
          - name: myapp-container
            image: busybox
            command: ['sh', '-c', 'echo Hello Kubernetes! I am updated!! && sleep 3600']
```

# Kubernetes Service

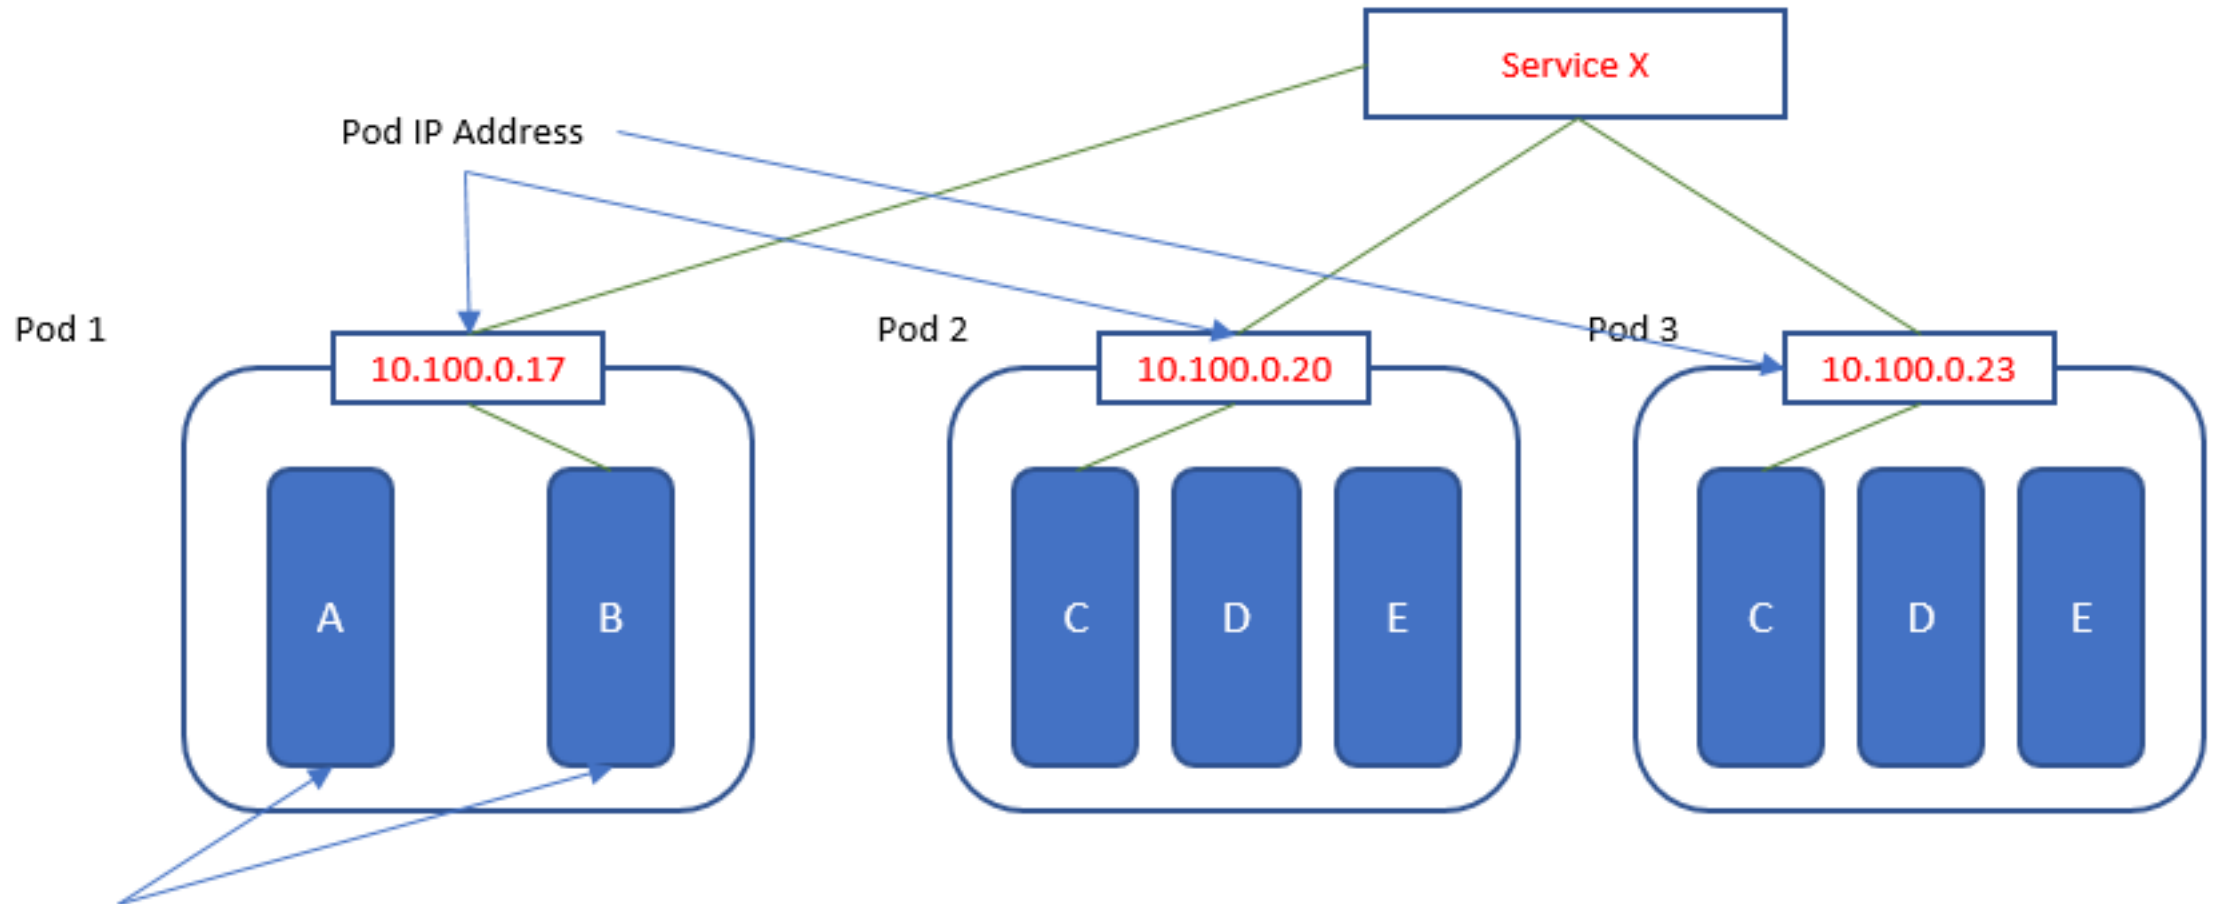
- Service
  - Exposes Pods to the outside as:
    - ClusterIP
    - NodePort
    - Load Balancer



# Service Example

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: hello-d
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

# Kubernetes Service

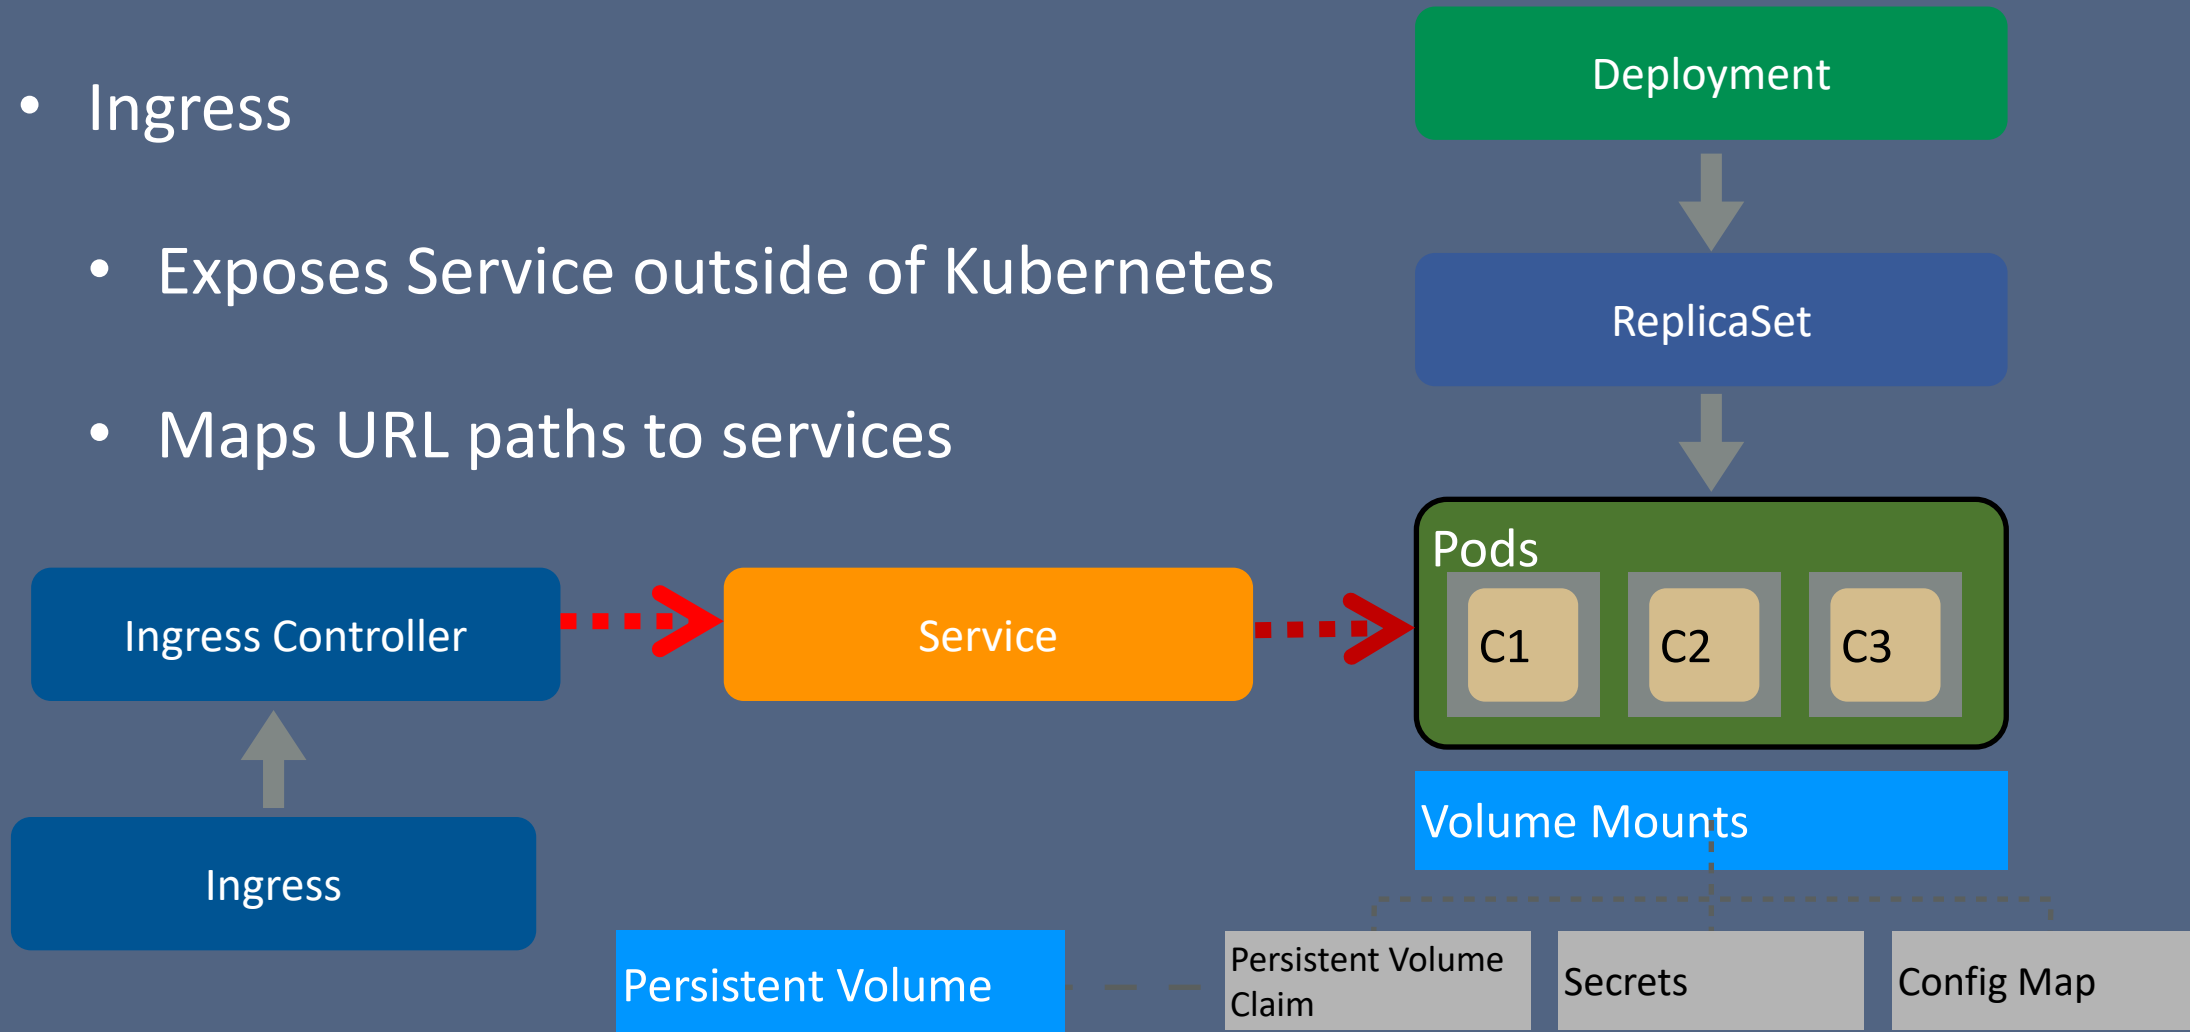


Containers

Container B accesses a function offered by container C (in either Pod 2 or 3) via a service

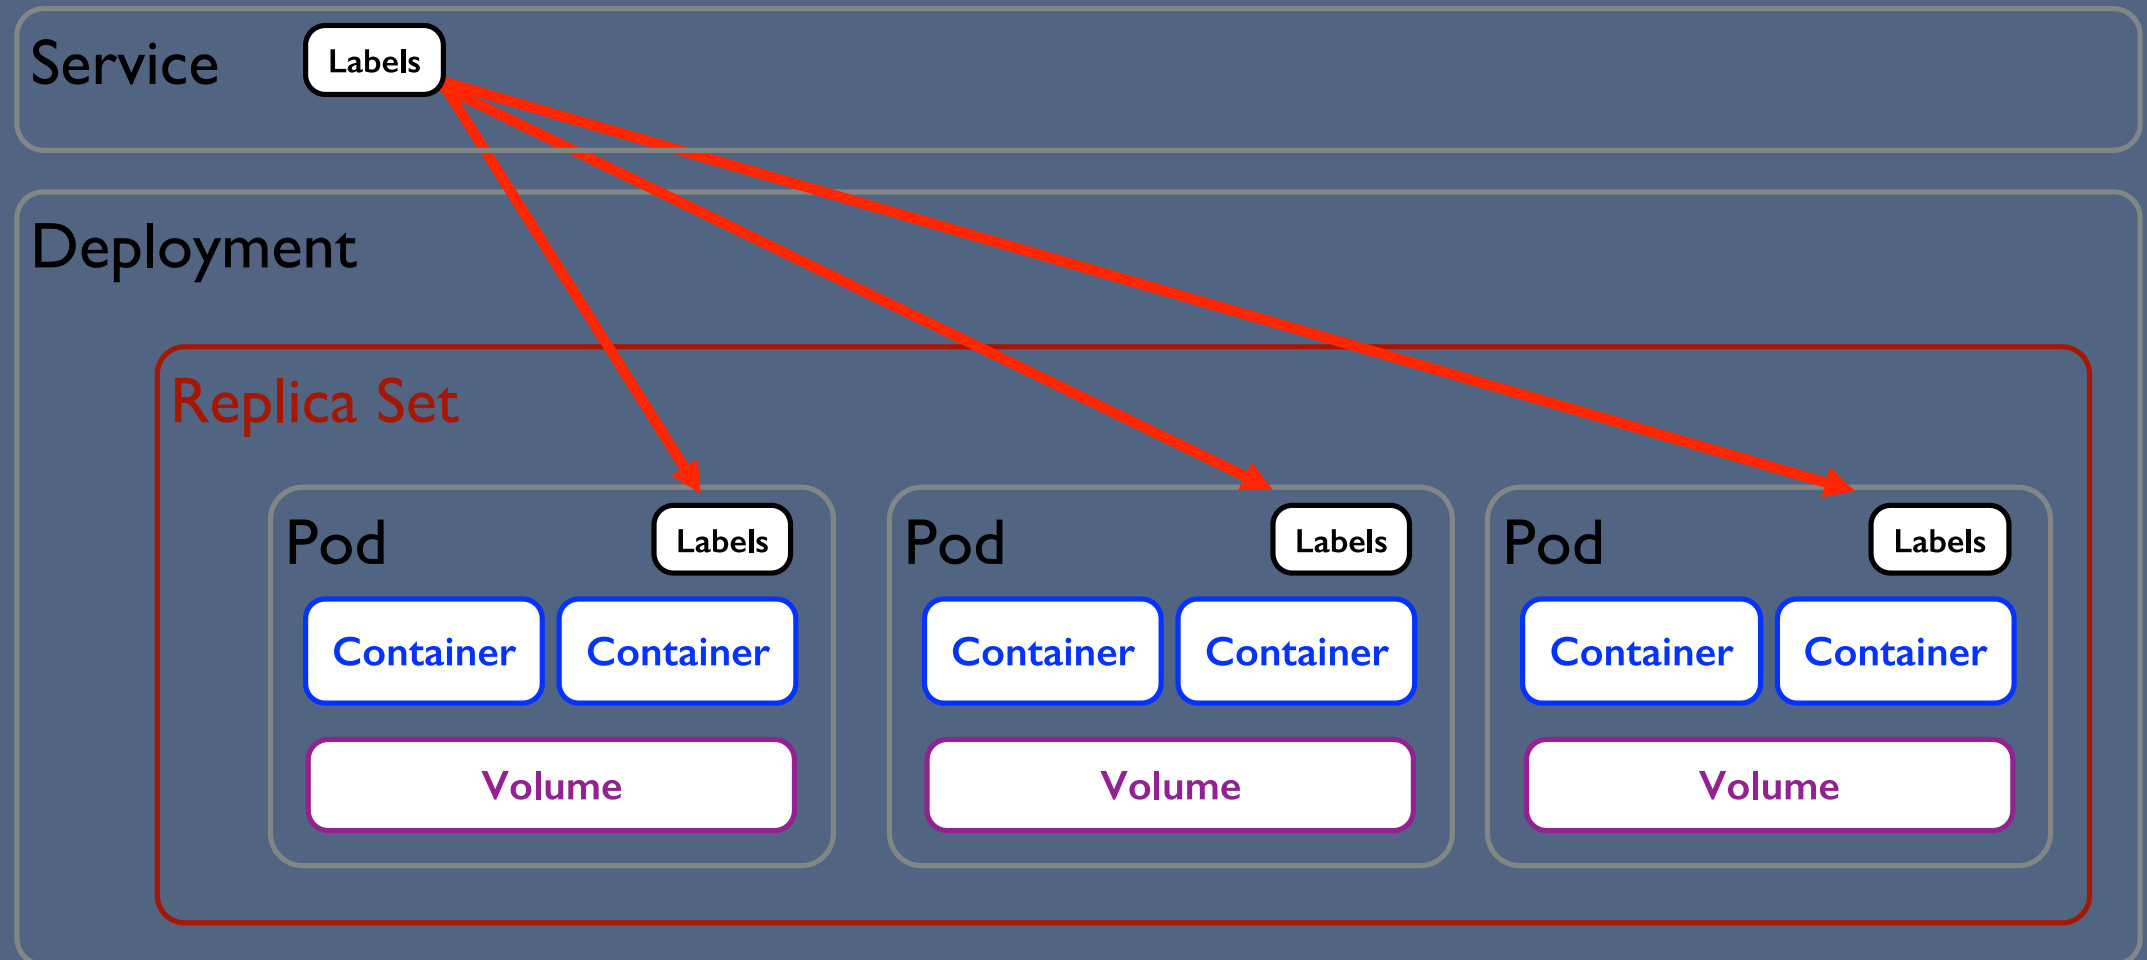
# Kubernetes Deployment Model

- Ingress
  - Exposes Service outside of Kubernetes
  - Maps URL paths to services



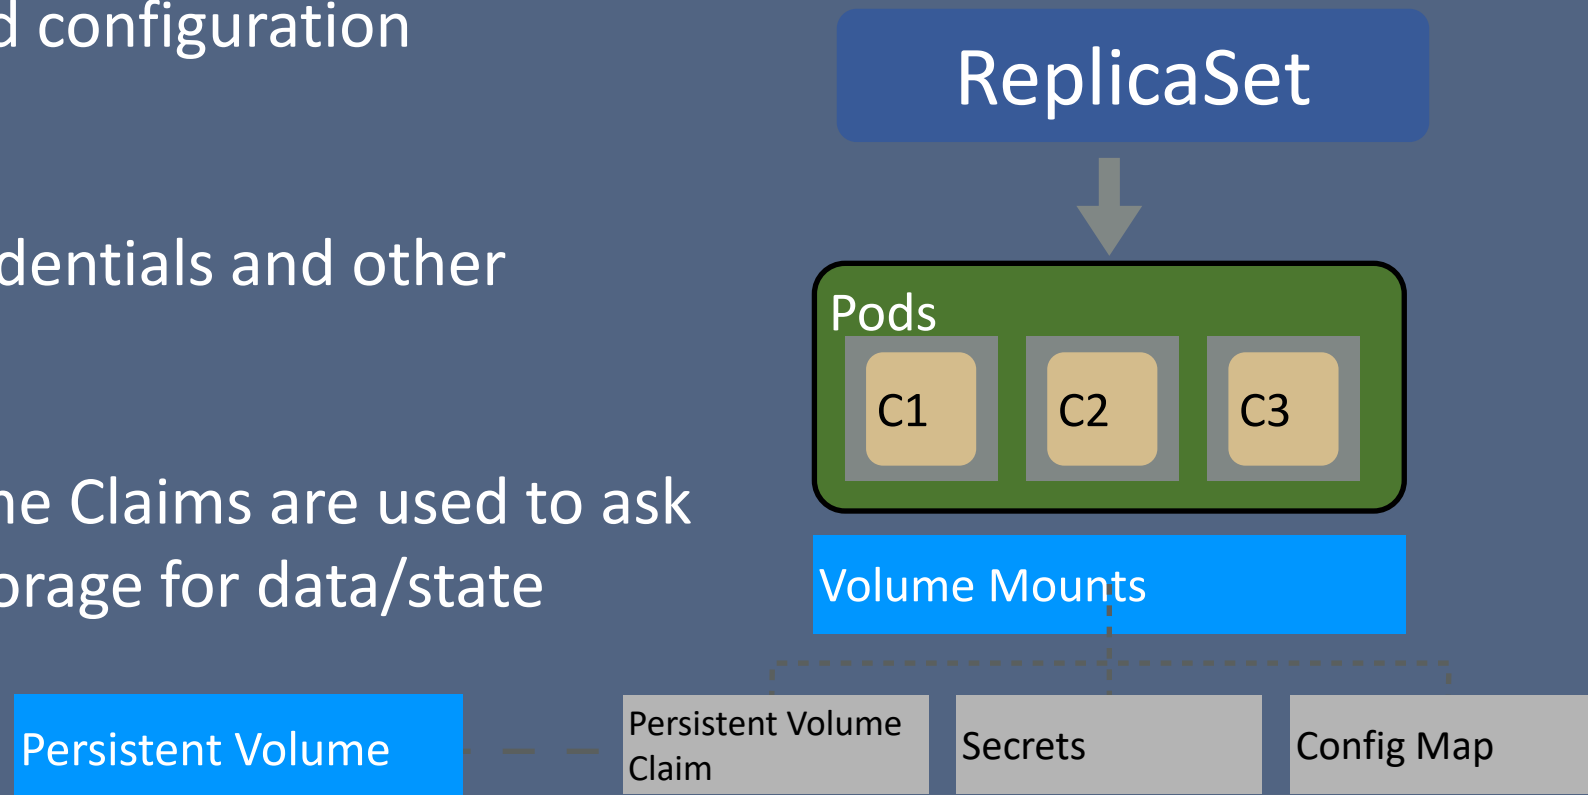


# Kubernetes Logical View: Putting all together




# Kubernetes Volume Mounts

- Volumes
  - ConfigMaps hold configuration parameters
  - Secrets hold credentials and other secrets
  - Persistent Volume Claims are used to ask for persistent storage for data/state



# Kubernetes machinery

- **Kube API Server**
- **Kube Scheduler**
- **Controller-manager**
- Kubelet
- Kube-proxy



Replacement



Plugin



Config

 kube-apiserver 

 kube-scheduler  



 kube-controller-manager 

 kubelet

 kube-proxy 

apiVersion: v1  
Kind: Deployment  
...

apiVersion: v1  
Kind: Service  
...



apiVersion: v1  
Kind: Deployment  
...

apiVersion: v1  
Kind: Service  
...




kube-apiserver



ETCD

apiVersion: v1  
Kind: Deployment  
...

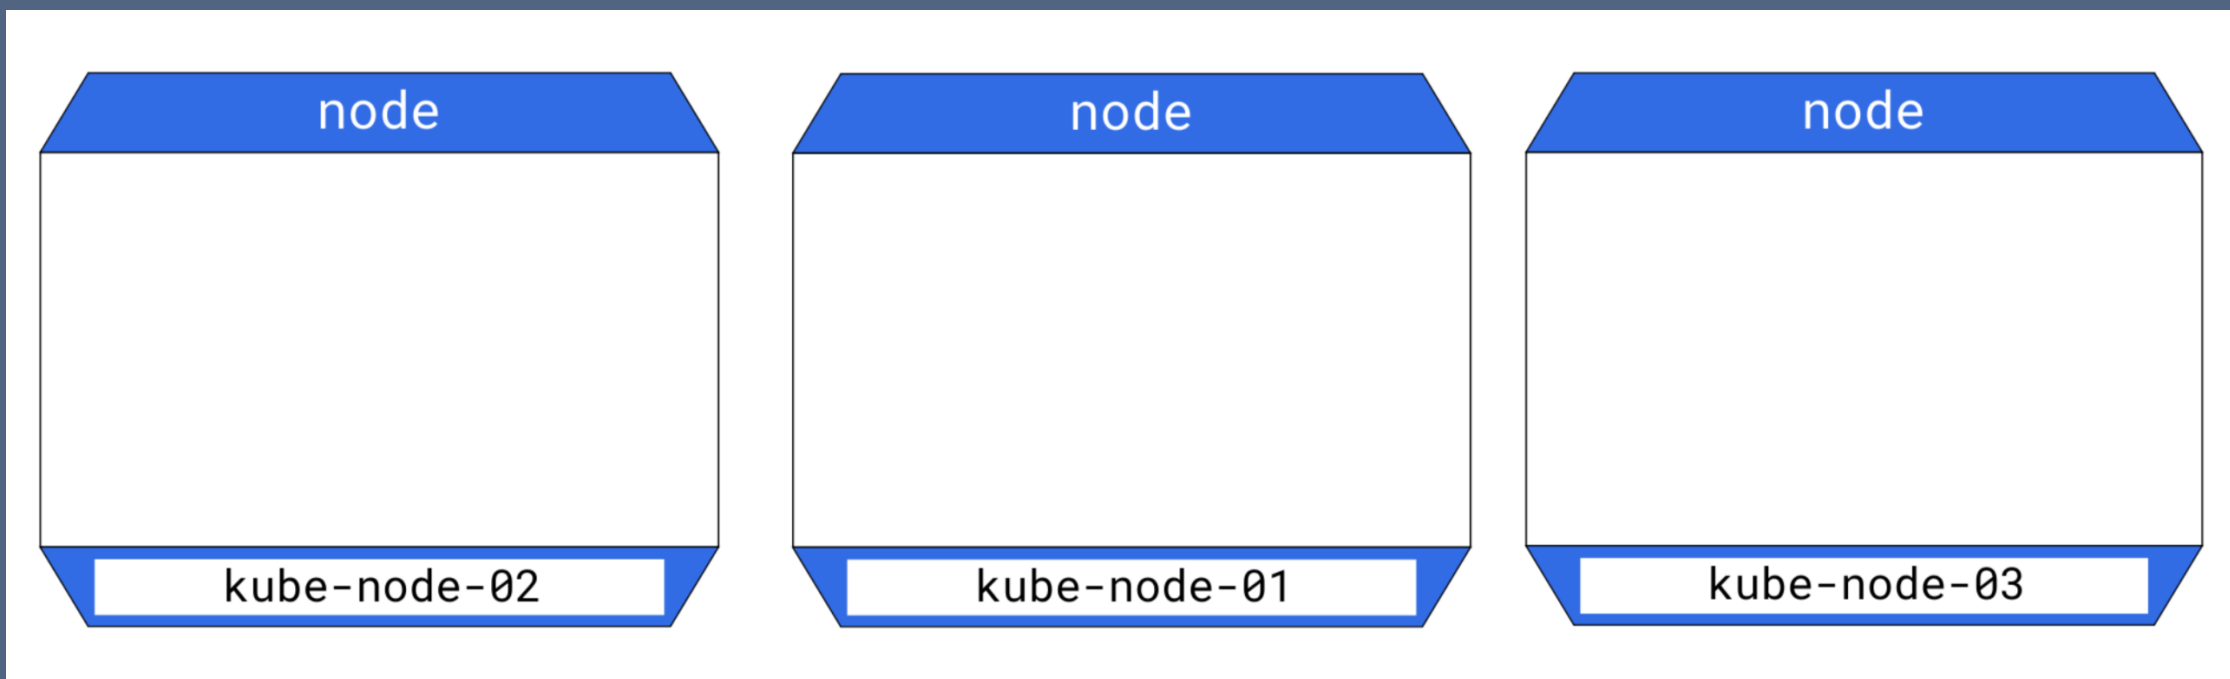

apiVersion: v1  
Kind: Service  
...

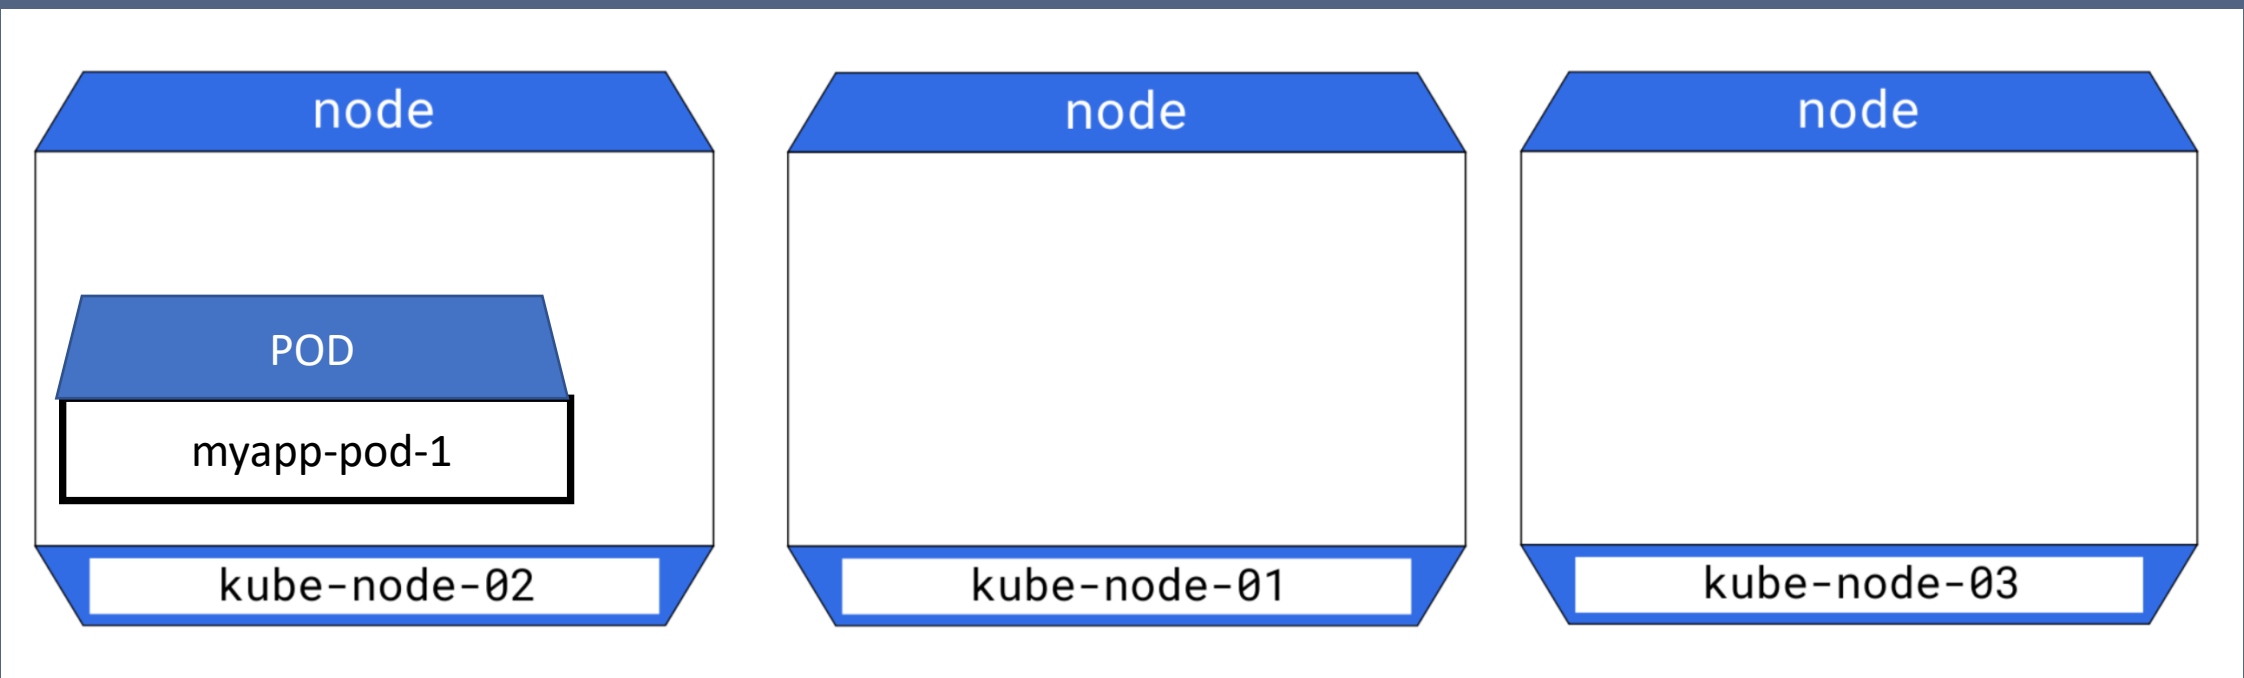








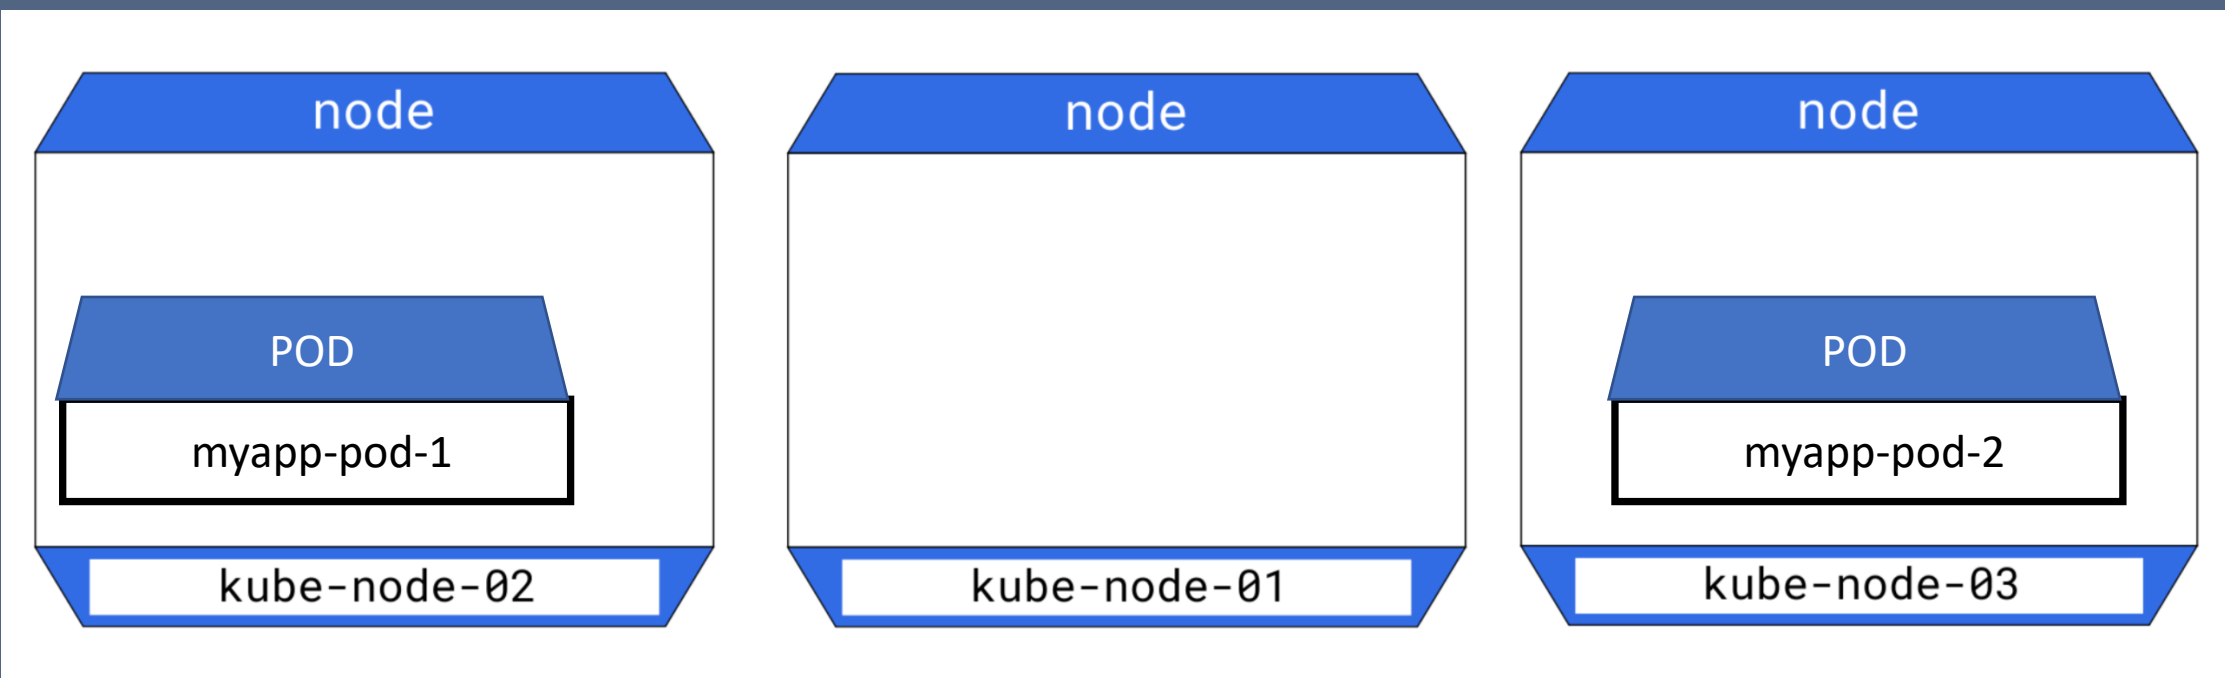


kube-scheduler












# kube-scheduler



affinity/anti-affinity  
nodeSelector  
taints/tolerations  
reservations/limits



custom  
schedulers

 kube-controller-manager



# kube-controller-manager

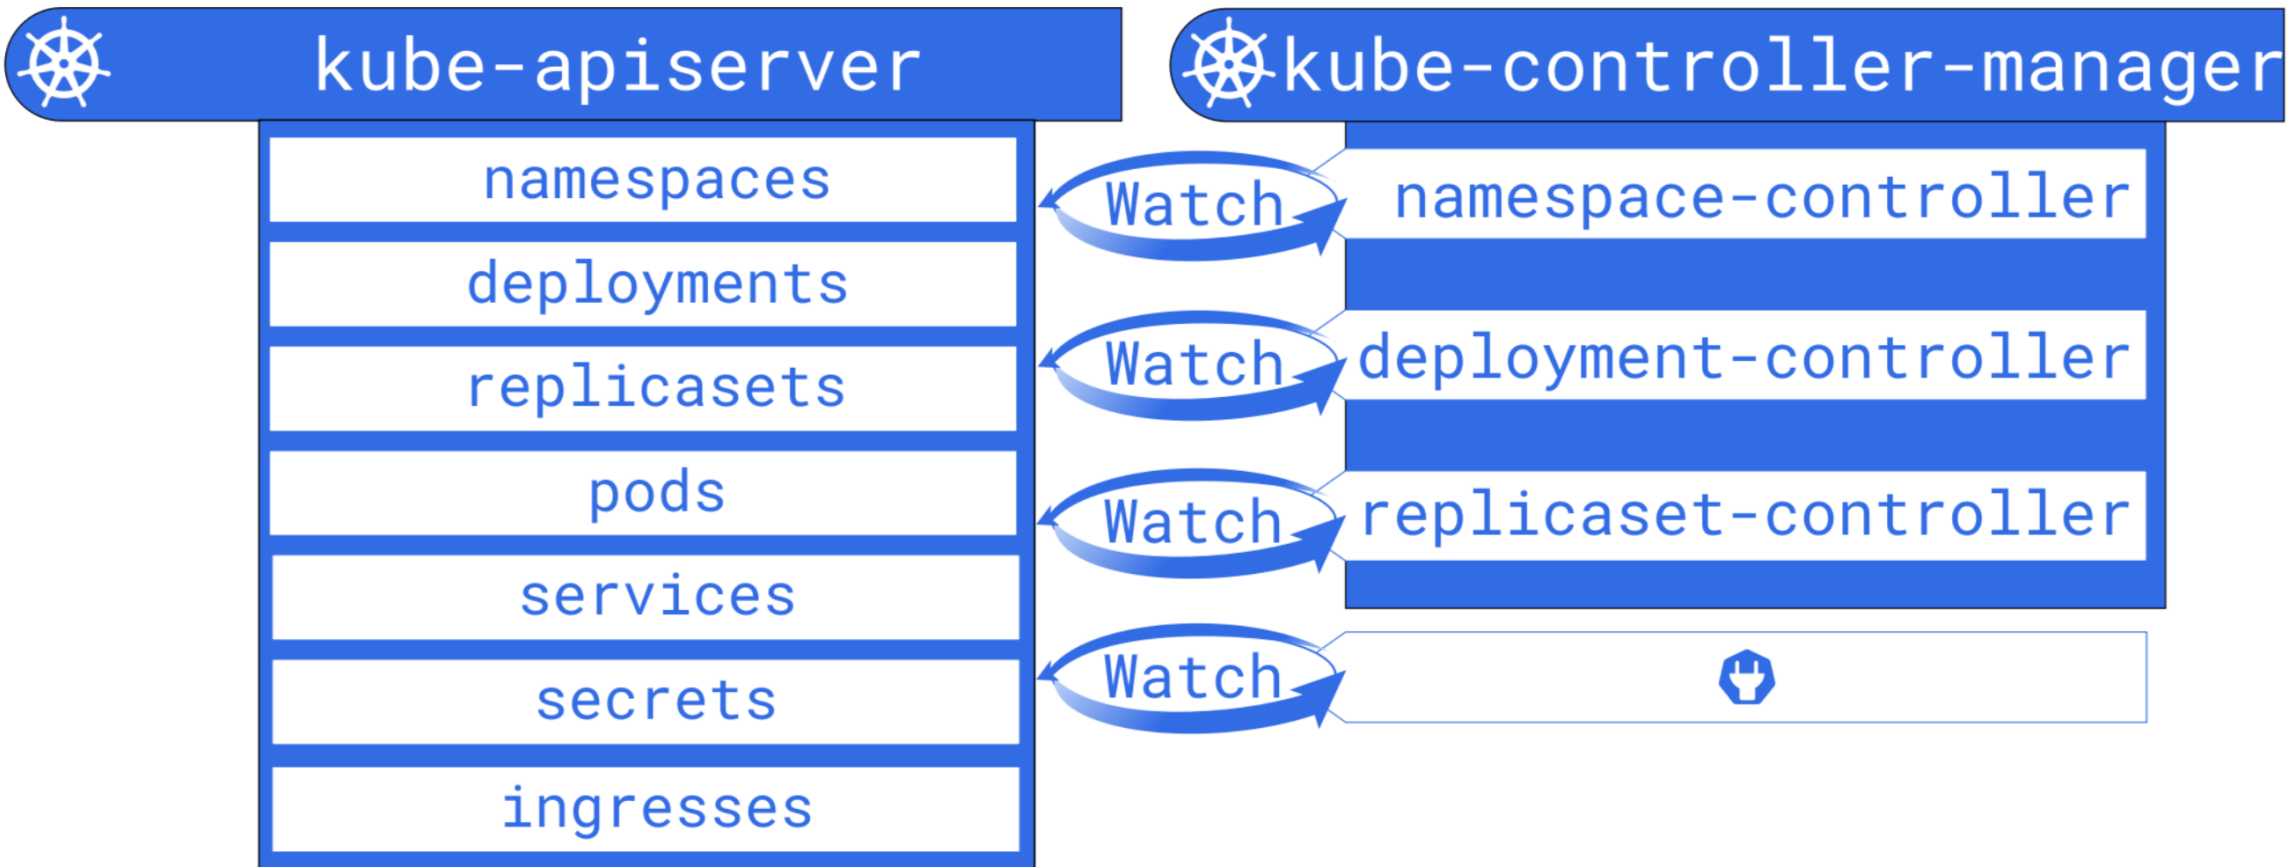
namespace-controller

deployment-controller

replicaset-controller










kubelet



# kubelet



## docker/rkt



node

POD

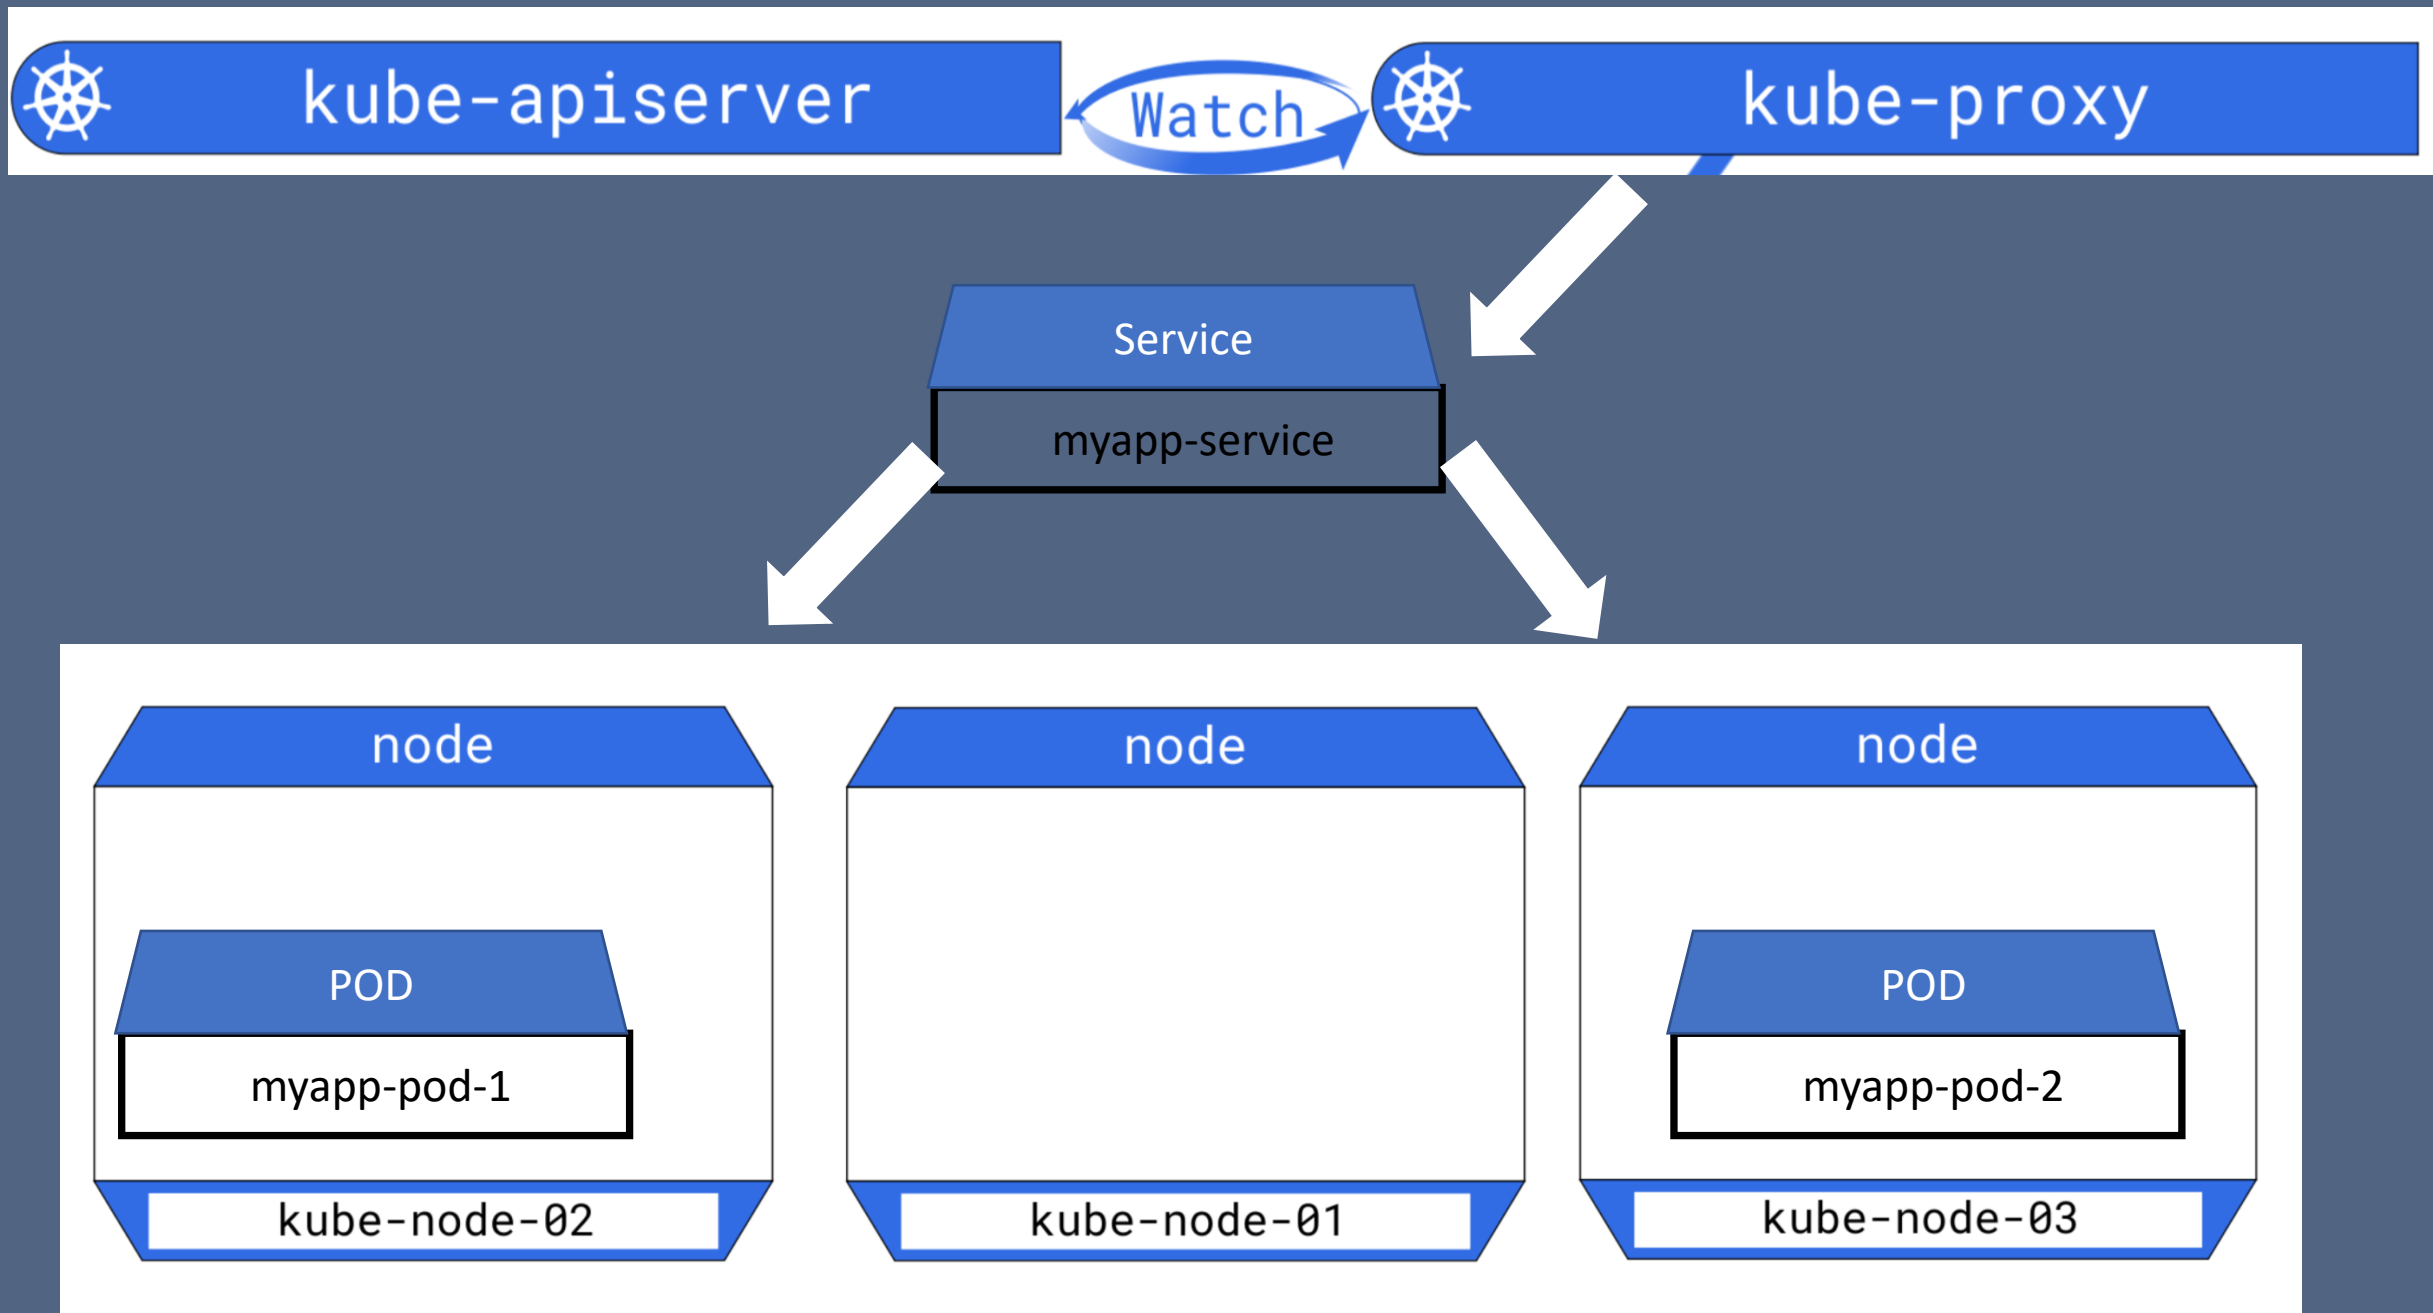
myapp-pod-1

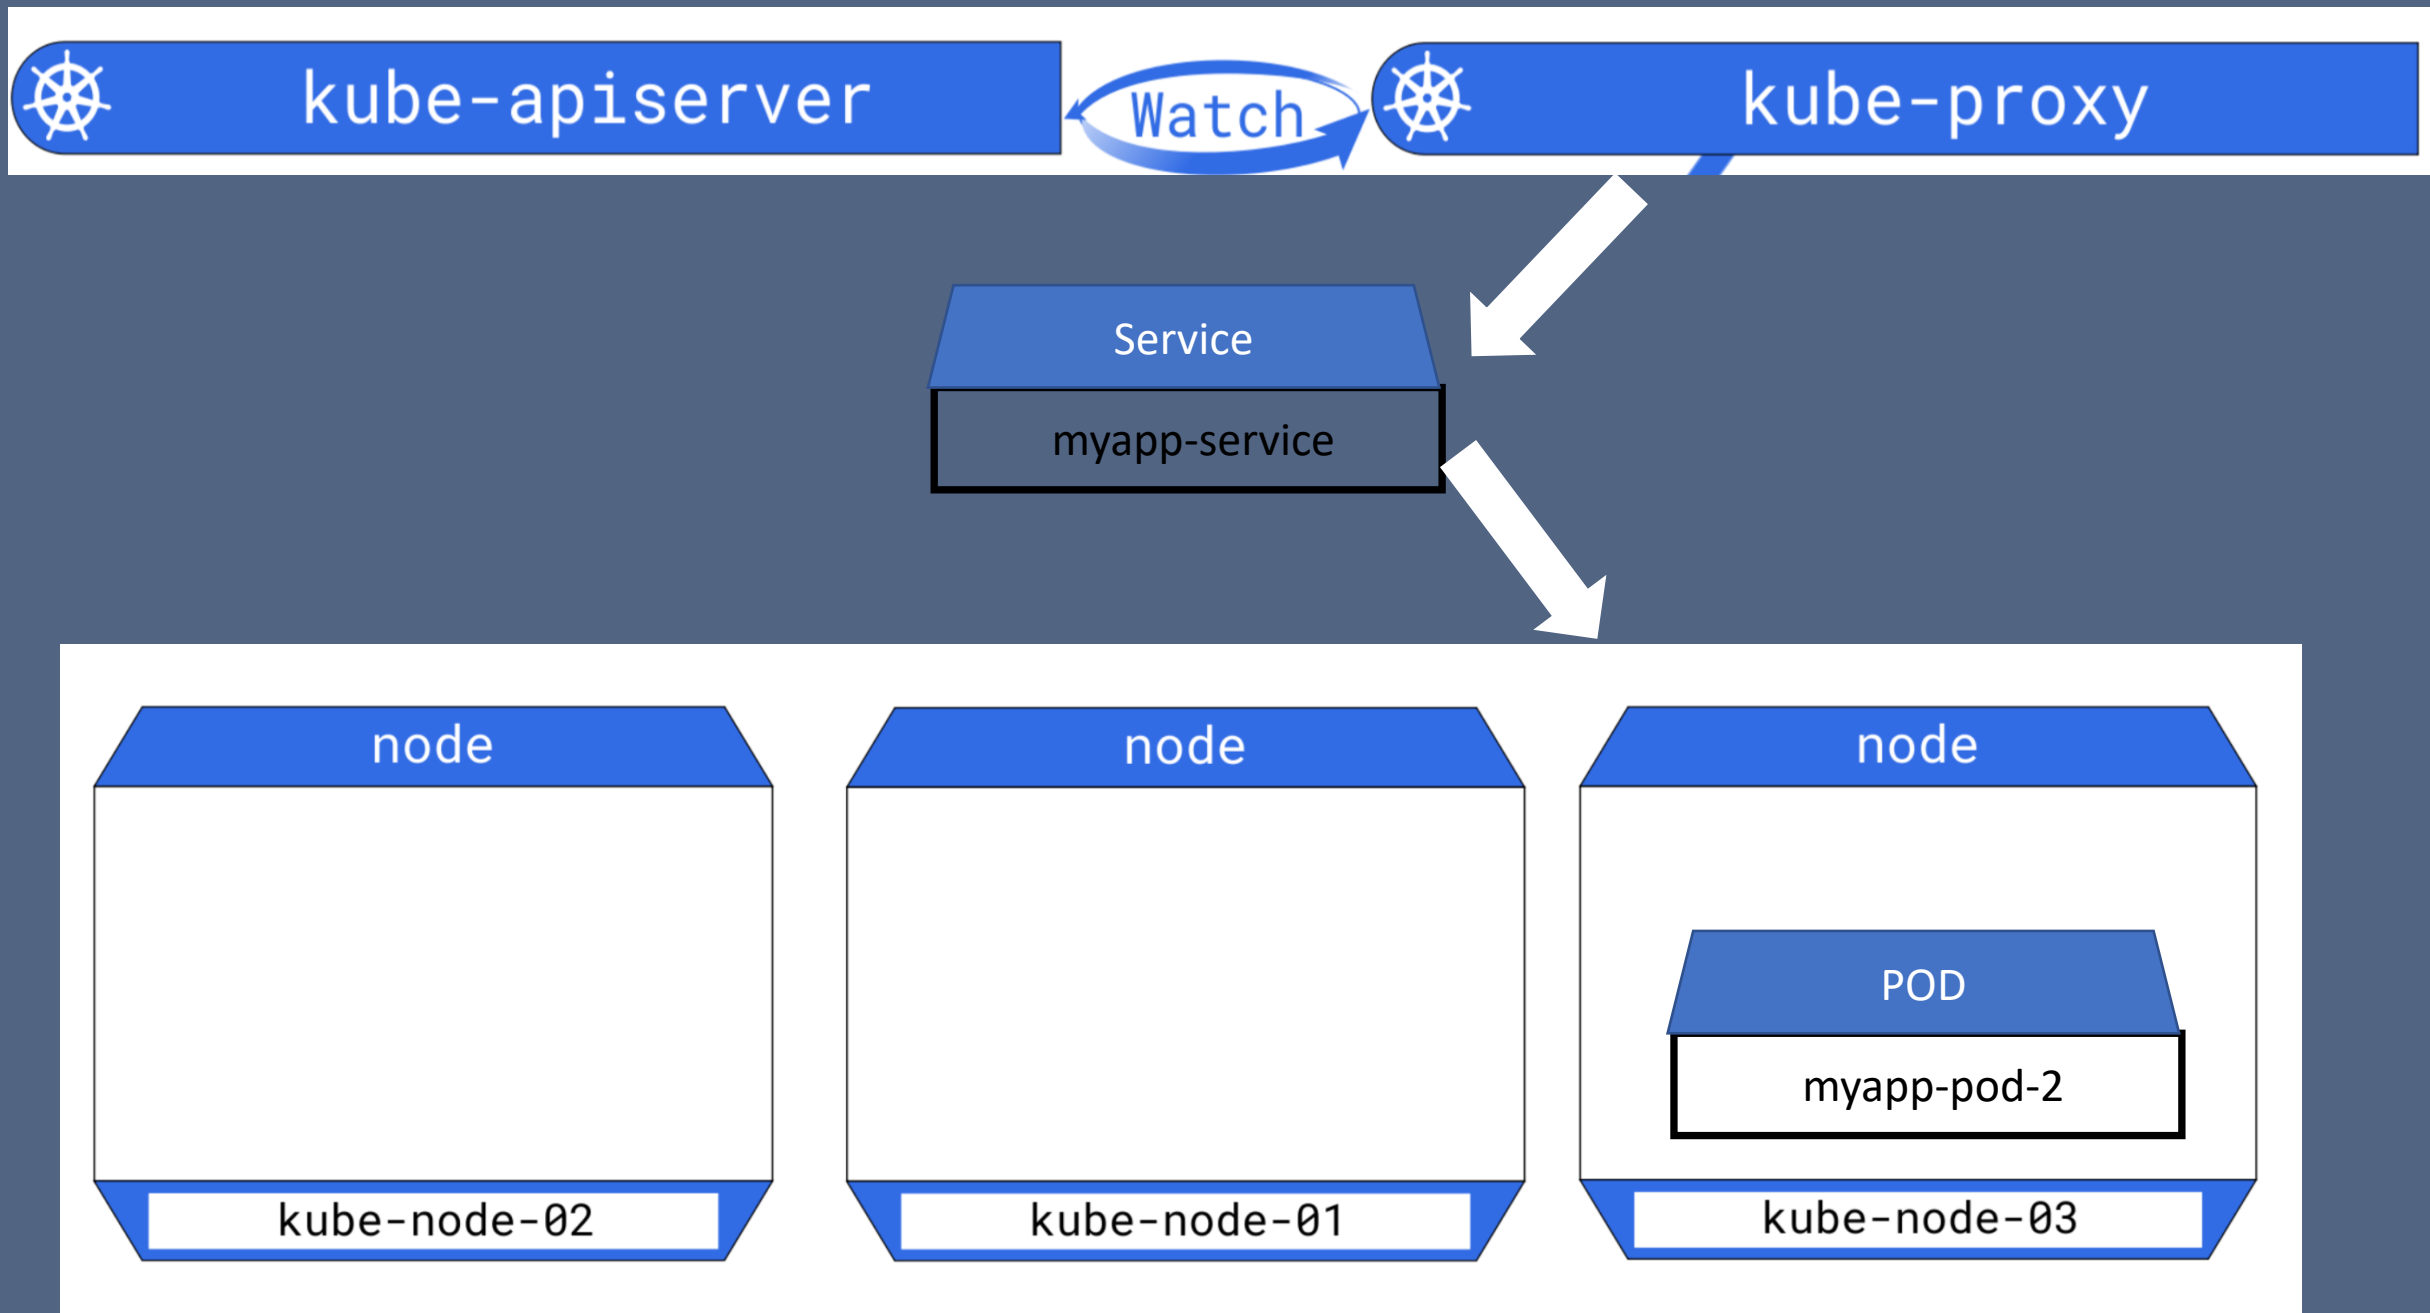
kube-node-01

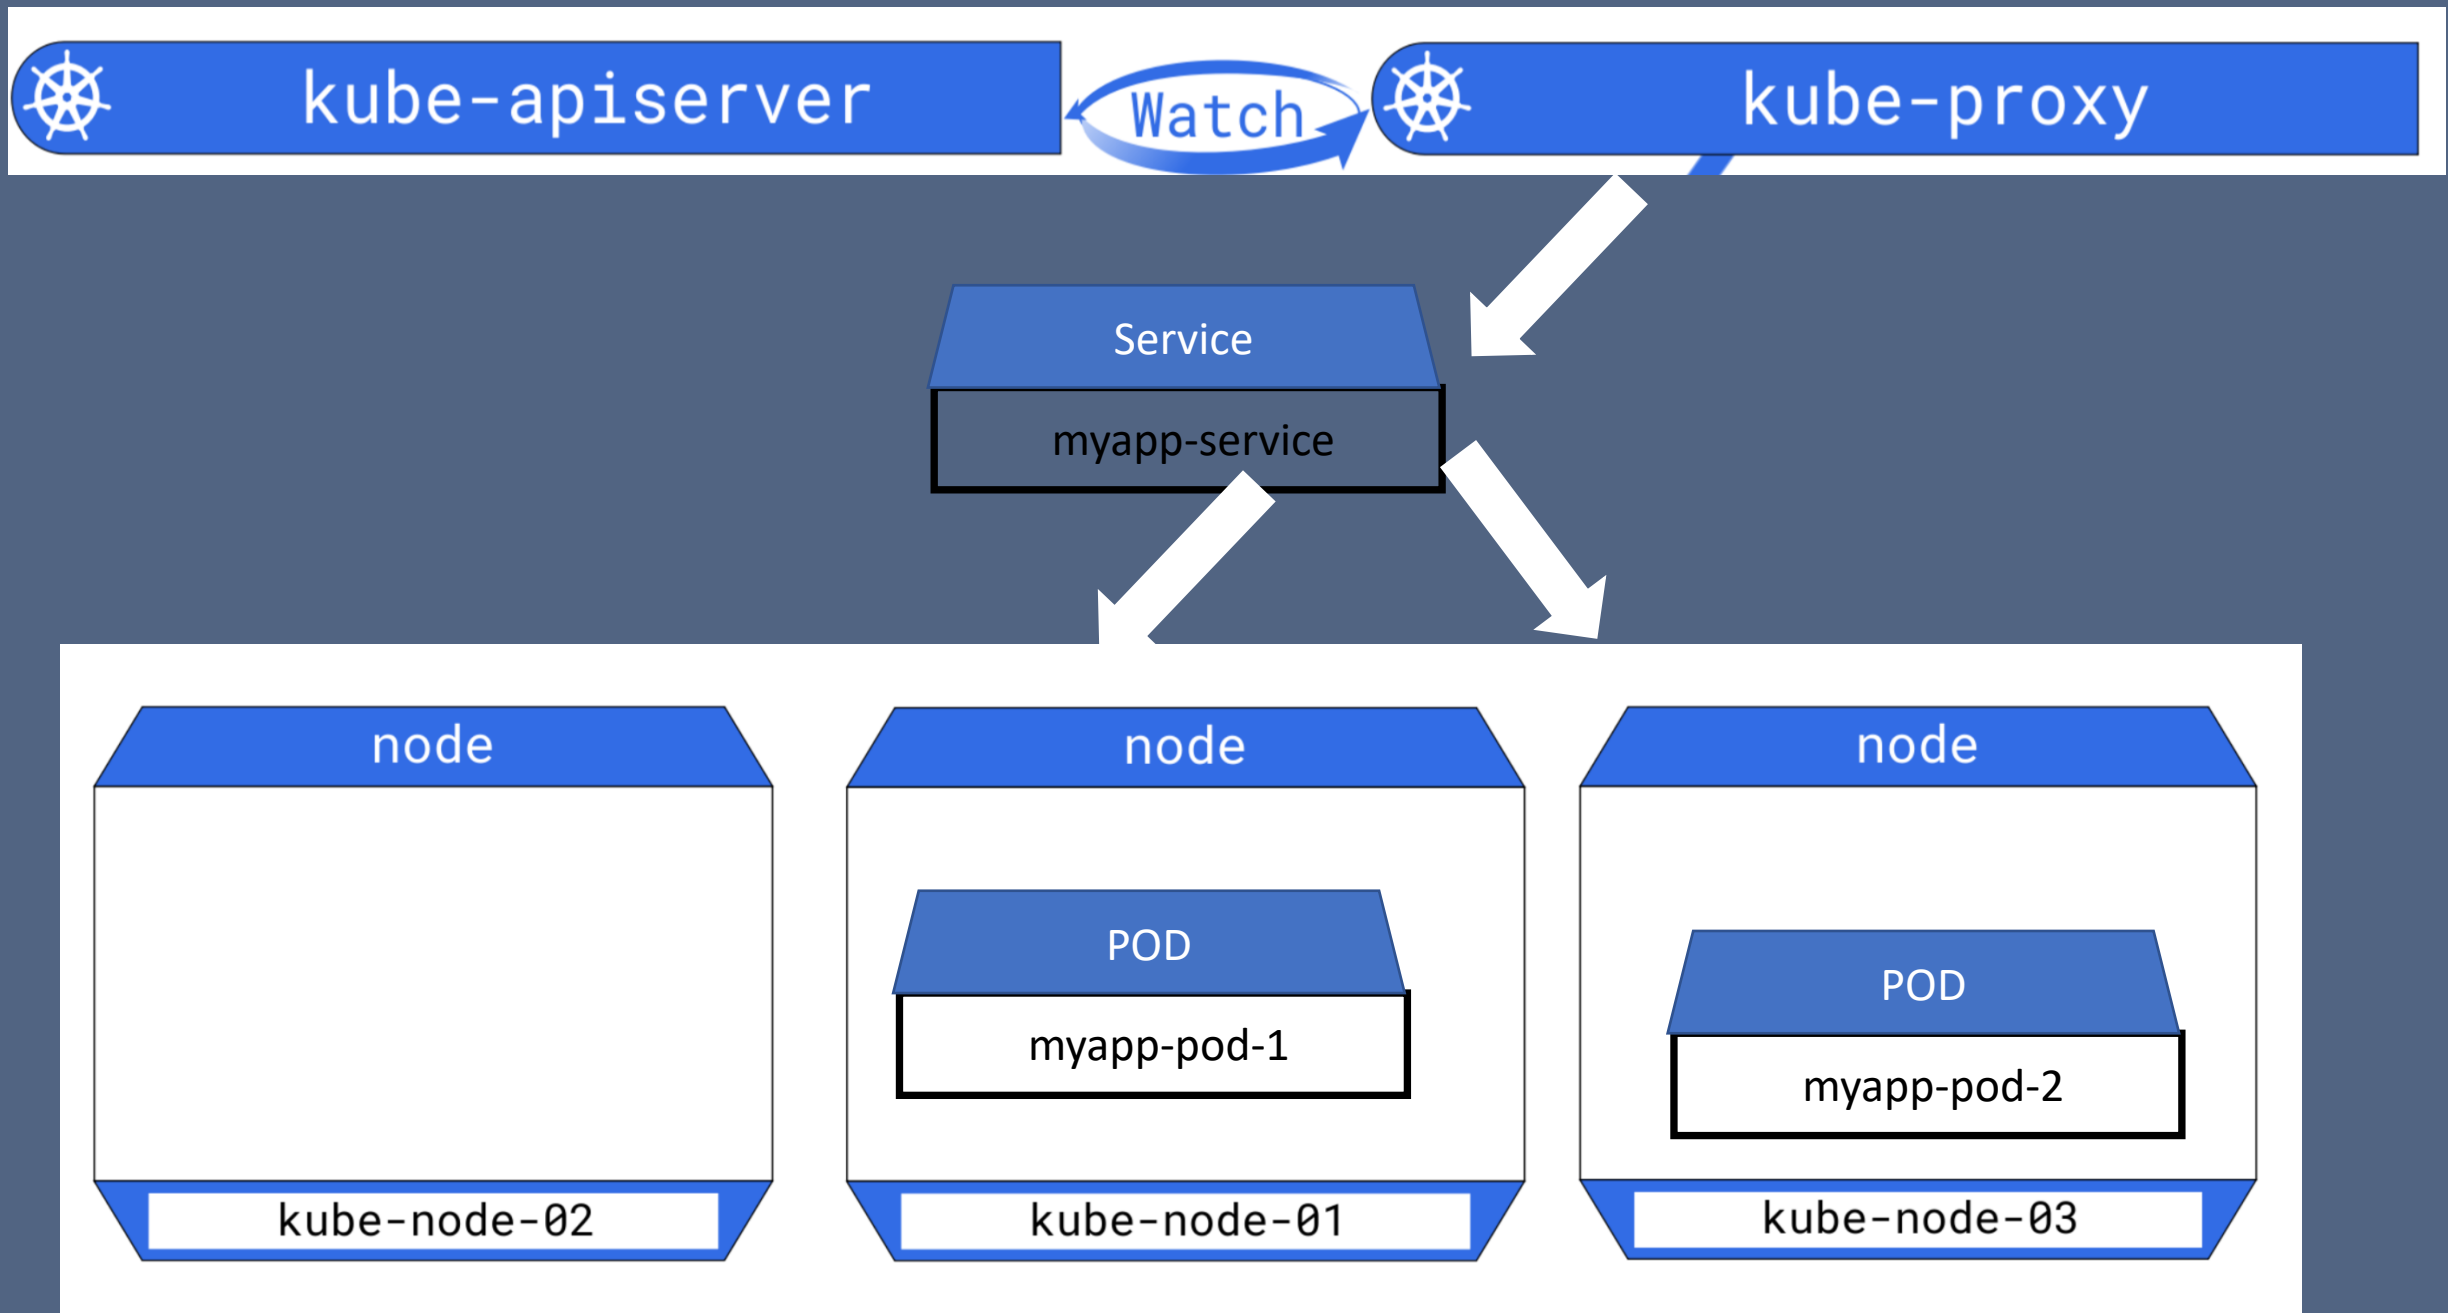




kube-proxy







# Lab

- Try the various example files here:
  - <https://github.com/seelam/k8s>



# Suggested Study Material

- Facebook Tupperware:
  - <https://engineering.fb.com/data-center-engineering/tupperware/>
- Kubernetes deconstructed
  - <http://kube-decon.carson-anderson.com/Layers/0-Intro.sozi.html#frame5378>
- The History of Kubernetes on a Timeline: <https://blog.risingstack.com/the-history-of-kubernetes/>
- The State of the Kubernetes Ecosystem
  - <https://thenewstack.io/ebooks/kubernetes/state-of-kubernetes-ecosystem/>